

# **ECON526: Quantitative Economics with Data Science Applications**

Applications of Linear Algebra

---

Jesse Perla

Overview

Difference Equations

Unemployment Dynamics

Latent Variables

Present Discounted Values

Discrete Latent Variables

(Optional) Matrix Conditioning and Stability

# Overview

---

# Motivation and Materials

- In this lecture, we will cover some applications of the tools we developed in the previous lecture
- The goal is to build some useful tools to sharpen your intuition on linear algebra and eigenvalues/eigenvectors, and practice some basic coding
- We introduce scikit-learn, a package for old-school (i.e. not deep learning or neural networks) ML and data analysis
  - Introduces “unsupervised learning” (i.e., tools to interpret data structure without any forecasts/predictions)
- Some additional material and references
  - QuantEcon Python
  - QuantEcon DataScience
  - A First Course in Quantitative Economics with Python

# Packages

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import scipy
4 from numpy.linalg import cond, matrix_rank, norm
5 from scipy.linalg import inv, solve, det, eig, lu, eigvals
6 from scipy.linalg import solve_triangular, eigvalsh, cholesky
```

# New Packages for Data Science and ML

```
1 import seaborn as sns
2 import pandas as pd
3 from sklearn.decomposition import PCA
4 from sklearn.cluster import KMeans
```

# Difference Equations

---



# Linear Difference Equations as Iterative Maps

- Consider  $A : \mathbb{R}^N \rightarrow \mathbb{R}^N$  as the linear map for the state  $x_t \in \mathbb{R}^N$
- An example of a linear difference equation is

$$x_{t+1} = Ax_t$$

where

$$A \equiv \begin{bmatrix} 0.9 & 0.1 \\ 0.5 & 0.8 \end{bmatrix}$$

```
1 A = np.array([[0.9, 0.1], [0.5, 0.8]])
2 x_0 = np.array([1, 1])
3 x_1 = A @ x_0
4 print(f"x_1 = {x_1}, x_2 = {A @ x_1}")
```

```
x_1 = [1.  1.3], x_2 = [1.03 1.54]
```

## Iterating with $\rho(A) > 1$

Iterate  $x_{t+1} = Ax_t$  from  $x_0$  for  $t = 100$

```
1 x_0 = np.array([1, 1])
2 t = 200
3 print(f"rho(A) = {np.max(np.abs(eigvals(A)))}")
4 print(f"x_{t} = {np.linalg.matrix_power(A, t) @ x_0}")
```

rho(A) = 1.079128784747792

x\_200 = [3406689.32410673 6102361.18640516]

- Diverges to  $x_\infty = \begin{bmatrix} \infty & \infty \end{bmatrix}^T$
- $\rho = 1 + 0.079$  says in the worst case (i.e.,  $x_t \propto$  the eigenvector associated with  $\lambda = 1.079$  eigenvalue), expands by 7.9% on each iteration

## Iterating with $\rho(A) < 1$

```
1 A = np.array([[0.6, 0.1], [0.5, 0.8]])
2 print(f"rho(A) = {np.max(np.abs(eigvals(A)))}")
3 print(f"x_{t} = {np.linalg.matrix_power(A, t) @ x_0}")
```

rho(A) = 0.9449489742783178

x\_200 = [6.03450418e-06 2.08159603e-05]

- Converges to  $x_\infty = \begin{bmatrix} 0 & 0 \end{bmatrix}^T$

## Iterating with $\rho(A) = 1$

- To make a matrix that has  $\rho(A) = 1$  reverse eigendecomposition!
- Leave previous eigenvectors in  $Q$ , change  $\Lambda$  to force  $\rho(A)$  directly

```
1 Q = np.array([[-0.85065081, -0.52573111], [0.52573111, -0.85065081]])
2 print(f"check orthogonal: dot(x_1,x_2) approx 0: {np.dot(Q[:,0], Q[:,1])}")
3 Lambda = [1.0, 0.8] # choosing eigenvalue so max_n|lambda_n| = 1
4 A = Q @ np.diag(Lambda) @ inv(Q)
5 print(f"rho(A) = {np.max(np.abs(eigvals(A)))}")
6 print(f"x_{t} = {np.linalg.matrix_power(A, t) @ x_0}")
```

```
check orthogonal: dot(x_1,x_2) approx 0: -1.9275984594779062e-17
```

```
rho(A) = 1.0
```

```
x_200 = [ 0.27639321 -0.17082039]
```

# Unemployment Dynamics

---

## Dynamics of Employment without Population Growth

- Consider an economy where in a given year  $\alpha = 5\%$  of employed workers lose job and  $\phi = 10\%$  of unemployed workers find a job
- We start with  $E_0 = 900,000$  employed workers,  $U_0 = 100,000$  unemployed workers, and no birth or death. Dynamics for the year:

$$E_{t+1} = (1 - \alpha)E_t + \phi U_t$$

$$U_{t+1} = \alpha E_t + (1 - \phi)U_t$$

- Can write this as a matrix equation

$$\underbrace{\begin{bmatrix} E_{t+1} \\ U_{t+1} \end{bmatrix}}_{X_{t+1}} = \underbrace{\begin{bmatrix} 1 - \alpha & \phi \\ \alpha & 1 - \phi \end{bmatrix}}_A \underbrace{\begin{bmatrix} E_t \\ U_t \end{bmatrix}}_{X_t}$$

# Simulating

Simulate by iterating  $X_{t+1} = AX_t$  from  $X_0$  until  $T = 100$

```
1 def simulate(A, X_0, T):
2     X = np.zeros((2, T+1))
3     X[:,0] = X_0
4     for t in range(T):
5         X[:,t+1] = A @ X[:,t]
6     return X
7 X_0 = np.array([900000, 100000])
8 A = np.array([[0.95, 0.1], [0.05, 0.9]])
9 T = 100
10 X = simulate(A, X_0, T)
11 print(f"X_{T} = {X[:,T]}")
```


X\_100 = [666666.6870779 333333.31292209]

## Plotting Code

```
1 fig, ax = plt.subplots(figsize=(6, 4))
2 ax.plot(range(T+1), X.T, label=["Employed", "Unemployed"])
3 ax.set(xlabel="t", ylabel="Number of Workers", title="Employment Status")
4 ax.legend()
5 plt.show()
```



# Dynamics of Unemployment



## Convergence to a Longrun Distribution

- Find  $X_\infty$  by iterating  $X_{t+1} = AX_t$  many times from a  $X_0$ ?
  - Check if it has converged with  $X_\infty \approx AX_\infty$
  - Is  $X_\infty$  the same from any  $X_0$ ? Will discuss “ergodicity” later
- Alternatively, note that this expression is the same as

$$1 \times \bar{X} = A\bar{X}$$

- i.e, a  $\lambda = 1$  where  $\bar{X}$  is the corresponding eigenvector of  $A$
- Is  $\lambda = 1$  always an eigenvalue? (yes if all  $\sum_{n=1}^N A_{ni} = 1$  for all  $i$ )
- Does  $\bar{X} = X_\infty$ ? For any  $X_0$ ?
- Multiple eigenvalues with  $\lambda = 1 \implies$  multiple  $\bar{X}$

## Using the First Eigenvector for the Steady State

```
1 Lambda, Q = eig(A)
2 print(f"real eigenvalues = {np.real(Lambda)}")
3 print(f"eigenvectors are column-by-column in Q =\n{Q}")
4 print(f"first eigenvalue = 1? {np.isclose(Lambda[0], 1.0)}")
5 X_bar = Q[:,0] / np.sum(Q[:,0]) * np.sum(X_0)
6 print(f"X_bar = {X_bar}\nX_{T} = {X[:,T]}")
```

```
real eigenvalues = [1.    0.85]
eigenvectors are column-by-column in Q =
[[ 0.89442719 -0.70710678]
 [ 0.4472136   0.70710678]]
first eigenvalue = 1? True
X_bar = [666666.66666667 333333.33333333]
X_100 = [666666.6870779  333333.31292209]
```


## Using the Second Eigenvalue for the Convergence Speed

- The second largest ( $\lambda_2 < 1$ ) provides information on the speed of convergence
  - 0 is instantaneous convergence here
  - 1 is no convergence here
- Create a new matrix with the same steady state, different speed

```
1 Lambda_fast = np.array([1.0, 0.5])
2 A_fast = Q @ np.diag(Lambda_fast) @ inv(Q) # same eigenvectors
3 X_fast = simulate(A_fast, X_0, T)
4 print(f"X_{T} = {X_fast[:,T]}")
```

```
X_100 = [666666.66666667 333333.33333334]
```

# Dynamics of Unemployment For Difference Convergence Speeds



# Latent Variables

---

# Features, Labels, and Latents

- Data science and ML often use different terminology than economists:
  - **Features** are economists **explanatory or independent variables**. They have the key source of variation to make predictions and conduct counterfactuals
  - **Labels** correspond to economists **observables or dependent variables**
  - **Latent Variables** are **unobserved variables**, typically sources of heterogeneity or which may drive both the dependent and independent variables
- Economists will use theory and experience to transform data (i.e., what ML people call “feature engineering”) for better explanatory power or map to theoretical models
- ML refers to methods using only **features** as **unsupervised learning**. The structure of the underlying data can teach you about its data generating process
- Key: uncover and interpret latent variables using statistics coupled with assumptions from economic theory. There is theory beyond all interpretation

# Principle Components and Factor Analysis

- Another application of eigenvalues is dimension reduction, which simplifies **features** by uncovering **latent** variables. Unsupervised
- One technique is Principle Components Analysis (PCA) which uncovers latent variables that capture the primary directions of variation in the underlying data
  - May allow mapping data into a lower-dimensional, uncorrelated set of features
  - Often uses Singular Value Decomposition (SVD) - a numerically stable generalization of eigendecomposition to non-square matrices. See QuantEcon SVD Notes
  - One of many methods. Many algorithms in ML and econometrics have similar goals but can be non-linear
- Given a matrix  $X \in \mathbb{R}^{N \times M}$ , can we find a lower-dimensional representation  $Z \in \mathbb{R}^{N \times L}$  for  $L < M$  that captures the most variation in  $X$ ?
- The columns of  $Z$  are called the principle components of  $X$
- The goal is to invert the  $X$  data to find the  $Z$ —and provide a mapping to reduce the dimensionality for future data



## Decomposing the Data

PCA typically uses SVD in practice - but we will use eigendecomposition (aka spectral decomposition if symmetric) instead

Start by doing a decomposition of the “covariance matrix” of the data,  $XX^T$ , and form diagonal  $\Lambda$  as a product of vectors  $\sigma \in \mathbb{R}^N$  (the singular values)

$$XX^T = Q\Lambda Q^T = \underbrace{Q\sigma}_X \underbrace{\sigma^T Q^T}_{(Q\sigma)^T = X^T}, \quad \text{where } \Lambda \equiv \sigma\sigma^T$$

Hence, denoting the  $n$ th column of  $Q$  as  $Q_n$ , we have

$$X = Q\sigma = Q_1\sigma_1 + Q_2\sigma_2 + \dots + Q_M\sigma_M$$

## Dimension Reduction

- Assume we sorted so  $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_M$ . Frequently  $\sigma_1 \gg \sigma_M$
- For many problems, the  $\sigma_m$  decay quickly, so we can approximate  $X$  with fewer terms by truncating the sum at  $L < M$ .

$$X \approx Q_1\sigma_1 + Q_2\sigma_2 + \dots + Q_L\sigma_L$$

- The eigendecomposition (or SVD) can find the orthogonal directions of the data that capture the most variation in the covariance matrix
  - Can prove it is the solution to the optimization problem to explain the most variation in the data with the lowest dimensionality
- This is useful even if it is not necessary to reduce the dimensionality of the data
  - Many high-dimensional data sources are low-dimensional in the suitable space.
  - This is especially true when models allow for nonlinear transformations (e.g., neural networks, autoencoders, etc.)

## Creating a Dataset with Latent Factors

Create a dataset with two latent factors, the first dominating

```
1 N = 50 # number of observations
2 L, M = 2, 3 # number of latent and observed factors
3 Z = np.random.randn(N, L) # latent factors
4 F = np.array([[1.0, 0.05], # X_1 = Z_1 + 0.05 Z_2
5               [2.0, 0.0], # X_2 = 2 Z_1
6               [3.0, 0.1]]) # X_3 = 3 Z_1 + 0.1 Z_2
7 X = Z @ F.T + 0.1 * np.random.randn(N, M) # added noise
```

## PCA without any Dimension Reduction

- See QuantEcon SVD for coding yourself. We will use the sklearn package
- The explained variance is the fraction of the variance explained by each factor

```
1  pca = PCA(n_components=3)
2  pca.fit(X)
3  with np.printoptions(precision=4, suppress=True, threshold=5):
4      print(f"Singular Values (sqrt eigenvalues):\n{pca.singular_values_}")
5      print(f"Explained Variance (ordered):\n{pca.explained_variance_ratio_}")
```

Singular Values (sqrt eigenvalues):

[29.4181 0.7984 0.6654]

Explained Variance (ordered):

[0.9988 0.0007 0.0005]

## Dimension Reduction with PCA

```
1  pca = PCA(n_components=2) # one less, and correctly specified
2  Z_hat = pca.fit_transform(X) # transformed by dropping last factor
3  # Scale and sign may not match due to indeterminacy
4  print(f"Correlation of Z_1 to Z_hat_1 = {np.corrcoef(Z.T, Z_hat.T)[0,2]}")
5  print(f"Correlation of Z_2 to Z_hat_2 = {np.corrcoef(Z.T, Z_hat.T)[1,3]}")
```

Correlation of Z\_1 to Z\_hat\_1 = -0.9995888625038096

Correlation of Z\_2 to Z\_hat\_2 = 0.5464398297674395

## Interpreting the Results

- The first factor in the decomposition is nearly perfectly (positive or negatively) correlated with the more important latent factor
  - The sign could have gone either way. The key is the shared information
  - How could you have known the sign is indeterminate?
- The 2nd factor has a good but not great correlation with the 2nd latent. Why?
- The variance decomposition that gave a 3rd factor with non-zero variance
  - In our process, there are only two latent variables. Why didn't it figure it out?
- How could you have changed the DGP to make this **less** successful?

# Present Discounted Values

---

## Geometric Series

- Assume dividends follow  $y_{t+1} = Gy_t$  for  $t = 0, 1, \dots$  and  $y_0$  is given
- $G > 0$ , dividends are discounted at factor  $\beta > 1$  then  $p_t = \sum_{s=0}^{\infty} \beta^s y_{t+s} = \frac{y_t}{1-\beta G}$
- More generally if  $x_{t+1} = Ax_t$ ,  $x_t \in \mathbb{R}^N$ ,  $y_t = Gx_t$  and  $A \in \mathbb{R}^{N \times N}$ , then

$$\begin{aligned} p_t &= y_t + \beta y_{t+1} + \beta^2 y_{t+2} + \dots = Gx_t + \beta GAx_t + \beta GA^2 x_t + \dots \\ &= \sum_{s=0}^{\infty} \beta^s A^s y_t \\ &= G(I - \beta A)^{-1} x_t \quad , \text{ if } \rho(A) < 1/\beta \end{aligned}$$

- i.e., spectral radius of  $A$ , the maximum scaling, must be less than discounting
- Intuition from univariate: of  $G \in \mathbb{R}^{1 \times 1}$  then  $\text{eig}(G) = G$ , so must have  $|\beta G| < 1$



## PDV Example

Here is an example with  $1 < \rho(A) < 1/\beta$ . Try with different  $A$

```
1 beta = 0.9
2 A = np.array([[0.85, 0.1], [0.2, 0.9]])
3 G = np.array([[1.0, 1.0]]) # row vector
4 x_0 = np.array([1.0, 1.0])
5 p_t = G @ solve(np.eye(2) - beta * A, x_0)
6 #p_t = G @ inv(np.eye(2) - beta * A) @ x_0 # alternative
7 rho_A = np.max(np.abs(np.real(eigvals(A))))
8 print(f"p_t = {p_t[0]:.4g}, spectral radius = {rho_A:.4g}, 1/beta = {1/beta:.4g}")
```

p\_t = 24.43, spectral radius = 1.019, 1/beta = 1.111


## A Portfolio Example

- Two assets pay dividends  $d_t \equiv [d_{1t} \ d_{2t}]^T$  following  $d_{t+1} = A d_t$  from  $d_0$
- Portfolio has  $G \equiv [G_1 \ G_2]$  shares of each asset and you discount at rate  $\beta$

```
1 A = np.array([[0.6619469, 0.49646018], [0.5840708, 0.4380531]])
2 G = np.array([[10.0, 4.0]])
3 d_0 = np.array([1.0, 1.0])
4 T, beta = 10, 0.9
5 p_0 = G @ solve(np.eye(2) - beta * A, d_0)
6 d = simulate(A, d_0, T)
7 y = G @ d # total dividends from portfolio
8 print(f"Portfolio value at t=0 is {p_0[0]:.5g}, total dividends at time {T}
```

Portfolio value at t=0 is 1424.5, total dividends at time 10 is 36.955

## Dividends Seem to Grow at a Similar Rate?



## Digging Deeper

- Let's do an eigendecomposition to analyze the factors

```
1 Lambda, Q = eig(A)
2 print(np.real(Lambda))
```

```
[ 1.10000000e+00 -2.65486732e-09]
```

- The first eigenvector is 1.1, but the second is (numerically) zero!
  - (In fact, I rigged it to be zero by constructing from a  $\Lambda$ , so this is all numerical copy/paste errors)
- Suggests that maybe only one latent factor driving both  $d_{1t}$  and  $d_{2t}$ ?

## Evolution Matrix is Very Simple with $\lambda_2 = 0$

If we stack columns  $Q \equiv \begin{bmatrix} q_1 & q_2 \end{bmatrix}$  then,

$$A = Q\Lambda Q^{-1} = Q \begin{bmatrix} \lambda_1 & 0 \\ 0 & 0 \end{bmatrix} Q^{-1} = \lambda_1 q_1 q_1^{-1}$$

```
1 lambda_1 = np.real(Lambda[0])  
2 q_1 = np.reshape(Q[:,0], (2,1))  
3 q_1_inv = np.reshape(inv(Q)[0,:], (1,2))  
4 norm(A - lambda_1 * q_1 @ q_1_inv) # pretty close to zero!
```

2.663274500543771e-09

# Transforming to the Latent State

- Recall:  $A = Q\Lambda Q^{-1}$  can be interpreted as:
  - Transformation to latent space, scaling, transform back
- We can demonstrate this in our example:
  - Transforming  $d_0$  to  $\ell_0$  using  $q_1^{-1}$
  - Evolving  $\ell_t$  from  $\ell_0$  with  $\ell_{t+1} = \lambda_1 \ell_t$ , or  $\ell_t = \lambda_1^t \ell_0$
  - Transforming back with  $q_1$
  - Checking if it aligns with the  $d_t$

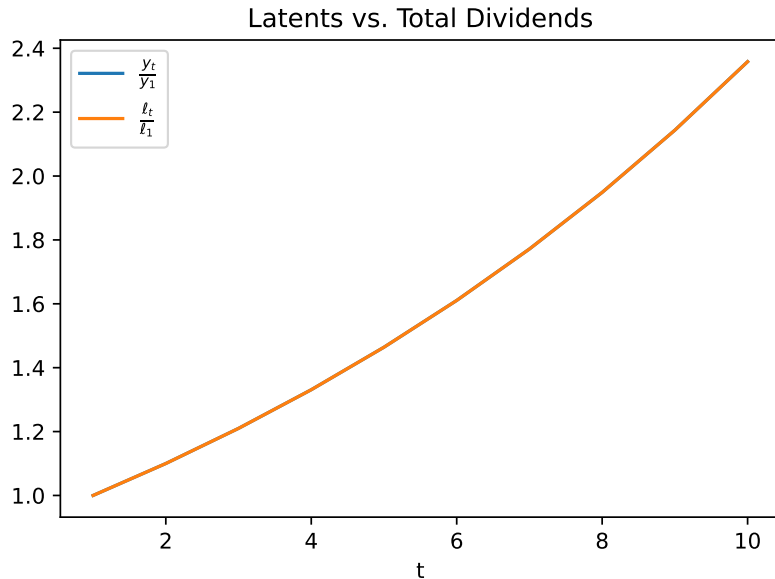
## Implementation

```
1 l_0 = lambda_1 * q_1_inv @ d_0 # latent space
2 l = l_0 * np.power(lambda_1, np.arange(0, T)) # powers
3 d_hat = q_1 * l # back to original space
4 # Missing d_0 since doing A * d_0 iterations
5 print(f"norm = {norm(d[:,1:] - d_hat)}")
6 y_hat = G @ d_hat
```

norm = 2.3494410877755447e-10

Let's see if these line up perfectly

# Total Dividends and the Latent Variable





# Discrete Latent Variables

---

# Clustering and Discrete Latent Variables

- PCA was a way to uncover continuous latent variables, or find low-dimensional continuous approximations
- But latent variables may be discrete (e.g., types of people, firms)
- Hidden discrete variables require assigning observations to groups
- Clustering lets you take a set of observations with (potentially) variables (i.e., features) and try to assign a discrete latent variable to each observation
  - Sometimes we know the number of groups from theory, usually we do not
  - While some are statistical and probabilistic, most methods assign a single latent type rather than a distribution
  - Choosing the number of groups to assign to is a challenge that requires theory and regularization - which we will avoid here
  - Instead, just as with PCA we will choose the number of groups ad-hoc rather than in a disciplined way

# Partitioning Sets

- Let  $X \in \mathbb{R}^{N \times M}$  with  $x_1, \dots, x_N \in \mathbb{R}^M$  the individual observations
- Assume that each  $x_n$  has a latent discrete  $k \in \{1, \dots, K\}$  then we can assign each observation to one group
  - $\mathbf{S} \equiv \{S_1, \dots, S_K\}$  where each  $n = 1, \dots, N$  is in exactly one  $S_k$  (i.e. a partition)
- The goal is to find the partition which is the most likely to assign each  $x_n$  the correct latent variable  $k$
- An alternative interpretation is to think of this as a dimension reduction technique which reduces complicated data into a low-dimensional discrete variable
- In economics we will sometimes cluster on some observations to reduce the dimension, then leave others continuous

## k-means Clustering

- If theory suggests that  $n \in S_k$  with similar latent variables should have similar  $x_n$ 
  - Group observations which are close or similar to each other
  - As always in linear algebra, close suggests using a norm. The euclidean norm in the  $M$  dimensional feature space is a good baseline
- The objective of k-means is to choose the partition  $\mathbf{S}$  which minimizes the norm between observations within each group (normalized by group size  $|S_k|$ ):

$$\min_{\mathbf{S}} \sum_{k=1}^K \frac{1}{|S_k|} \sum_{x_n, x_{n'} \in S_k} \|x_n - x_{n'}\|_2^2$$

- Using standard euclidean norm between two elements in  $S_k$

$$\|x_n - x_{n'}\|_2^2 = \sum_{m=1}^M (x_{nm} - x_{n'm})^2$$

## k-means Objective Function

- Can prove that the previous objective is equivalent to minimizing the sum of the squared distances from the group  $k$ 's mean

$$\min_{\mathbf{S}} \sum_{k=1}^K \sum_{n \in S_k} \|x_n - \bar{x}_k\|_2^2$$

- Where the mean of group  $k$  is standard, and across all  $m$  features

$$\bar{x}_k \equiv \frac{1}{|S_k|} \sum_{x_n \in S_k} x_n$$

- Careful with using wildly different scales (i.e.  $\bar{x}_k$  may be dominated by one feature)

## Generating Data with Latent Groups


Generate data with 2 features and 2 latent groups and see how k-means does

```
1 mu_1 = np.array([0.0, 0.0]) # mean of k=1
2 mu_2 = np.array([1.0, 1.0]) # mean of k=2
3 sigma = np.array([[0.2, 0], [0, 0.2]]) # use same variance
4 N = 100 # observations
5 X_1 = np.random.multivariate_normal(mu_1, sigma, N)
6 X_2 = np.random.multivariate_normal(mu_2, sigma, N)
7 df_1 = pd.DataFrame({"f1": X_1[:, 0], "f2": X_1[:, 1], "k": 1})
8 df_2 = pd.DataFrame({"f1": X_2[:, 0], "f2": X_2[:, 1], "k": 2})
9 df = pd.concat([df_1, df_2], ignore_index=True)
```

## Plotting Code with Seaborn

```
1 fig, ax = plt.subplots(figsize=(6, 4))
2 sns.scatterplot(data=df, x="f1", y="f2", hue="k", ax=ax)
3 ax.set(xlabel="Feature 1", ylabel="Feature 2", title="Latent Groups")
4 plt.show()
```

## Plot of Features and Latents





## k-means to Recover the Latent Groups

- Run k-means with 2 clusters and check the results
- If correlation is close to 1 then successfully recovered the latent groups
- If the correlation is close to -1 then it was succesful. The latent groups  $\hat{k}$  numbers are ordered arbitrarily, just as  $k$  was

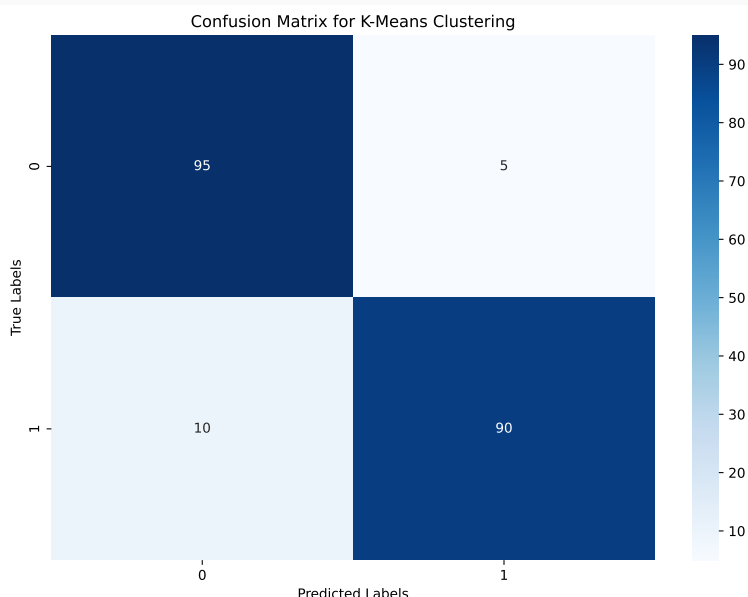
```
1 kmeans = KMeans(n_clusters=2, random_state=0)
2 k_hat = kmeans.fit_predict(df[["f1", "f2"]])
3 df["k_hat"] = k_hat + 1
4 corr = df["k"].corr(df["k_hat"])
5 print(f"Correlation between k and k_hat:{corr:.2f}")
```

Correlation between k and k\_hat:0.85

# Confusion Matrix

```
1 from sklearn.metrics import confusion_matrix
2
3 # compute confusion matrix
4 cm = confusion_matrix(df["k"], df["k_hat"])
5
6 # plot confusion matrix
7 sns.heatmap(cm, annot=True, cmap='Blues')
8 plt.xlabel('Predicted k')
9 plt.ylabel('True k')
10 plt.title('Confusion Matrix for K-Means Clustering')
11 plt.show()
```

# Confusion Matrix



## Swap $\hat{k}$ and Compare


```
1 df['k_hat'] = df['k_hat'].replace({1: 2, 2: 1}) # swap for this example
2 print(f"Correlation now {df['k'].corr(df['k_hat'])}")
3 df['k_cluster'] = df.apply(lambda x: x['k'] if x['k'] == x['k_hat'] else 3
4 # Add text to display using latex
5 df['Cluster'] = df['k_cluster'].replace({1: r'$k=\hat{k}=1$',
6                                           2: r'$k=\hat{k}=2$',
7                                           3: r'$k \neq \hat{k}$'})
```

Correlation now -0.851064496346989

## Plotting the Uncovered Latent Groups

```
1 fig, ax = plt.subplots(figsize=(6, 4))
2 sns.scatterplot(data=df, x="f1", y="f2", hue="Cluster", ax=ax)
3 ax.set(xlabel="Feature 1", ylabel="Feature 2", title="k-means Recovered Gr
4 plt.show()
```

## Plotting the Uncovered Latent Groups



## **(Optional) Matrix Conditioning and Stability**

---

# Matrix Conditioning

- Poorly conditioned matrices can lead to inaccurate or wrong solutions
- Tends to happen when matrices are close to singular or when they have very different scales - so there will be times when you need to rescale your problems

```
1 eps = 1e-7
2 A = np.array([[1, 1], [1 + eps, 1]])
3 print(f"A =\n{A}")
4 print(f"A^{-1} =\n{inv(A)}")
```

```
A =
[[1.          1.          ]
 [1.00000001  1.          ]]
A^{-1} =
[[-9999999.99336215  9999999.99336215]
 [10000000.99336215 -9999999.99336215]]
```



## Condition Numbers of Matrices

- $\det(A) \approx 0$  may say it is “almost” singular, but it is not scale-invariant
- $\text{cond}(A) \equiv \|A\| \cdot \|A^{-1}\|$  where  $\|\cdot\|$  is the matrix norm - expensive to calculate in practice. Connected to eigenvalues  $\text{cond}(A) = \left| \frac{\lambda_{\max}}{\lambda_{\min}} \right|$
- Scale free measure of numerical issues for a variety of matrix operations
- Intuition: if  $\text{cond}(A) = K$ , then  $b \rightarrow b + \nabla b$  change in  $b$  amplifies to a  $x \rightarrow x + K\nabla b$  error when solving  $Ax = b$ .
- See Matlab Docs on `inv` for example, where `inv` is a bad idea due to poor conditioning

```
1 print(f"condition(I) = {cond(np.eye(2))}")  
2 print(f"condition(A) = {cond(A)}, condition(A^(-1)) = {cond(inv(A))}")
```

```
condition(I) = 1.0
```

```
condition(A) = 40000001.939191714, condition(A^(-1)) = 40000002.00307444
```

## Example with Interpolation

- Consider fitting data  $x \in \mathbb{R}^{N+1}$  and  $y \in \mathbb{R}^{N+1}$  with an  $N$ -degree polynomial
- That is, find  $c \in \mathbb{R}^{N+1}$  such that

$$c_0 + c_1x_1 + c_2x_1^2 + \dots + c_Nx_1^N = y_1$$

$$\dots = \dots$$

$$c_0 + c_1x_N + c_2x_N^2 + \dots + c_Nx_N^N = y_N$$

- Which we can then use as  $P(x) = \sum_{n=0}^N c_nx^n$  to interpolate between the points

## Writing as a Linear System

- Define a matrix of all of the powers of the  $x$  values

$$A \equiv \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^N \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^N \end{bmatrix}$$

- Then solve for  $c$  as the solution to,

$$Ac = y$$

- Which we can solve using our tools. As long as  $x_n$  are unique, it is  $A$  is invertible
- Let's look at the numerical error here from the interpolation using the inf-norm, i.e.,  $\|x\|_\infty = \max_n |x_n|$

## Solving an Example

```
1 N = 5
2 x = np.linspace(0.0, 10.0, N + 1)
3 y = np.exp(x) # example function to interpolate
4 A = np.array([[x_i**n for n in range(N + 1)] for x_i in x]) # or np.vander
5 c = solve(A, y)
6 c_inv = inv(A) @ y
7 print(f"error = {norm(A @ c - y, np.inf)}, \
8 error using inv(A) = {norm(A @ c_inv - y, np.inf)}")
9 print(f"cond(A) = {cond(A)}")
```

```
error = 2.2737367544323206e-11, error using inv(A) = 1.0986695997416973e-09
cond(A) = 564652.321404467
```

## Things Getting Poorly Conditioned Quickly

```
1 N = 10
2 x = np.linspace(0.0, 10.0, N + 1)
3 y = np.exp(x) # example function to interpolate
4 A = np.array([[x_i**n for n in range(N + 1)] for x_i in x]) # or np.vander
5 c = solve(A, y)
6 c_inv = inv(A) @ y # Solving with inv(A) instead of solve(A, y)
7 print(f"error = {norm(A @ c - y, np.inf)}, \
8 error using inv(A) = {norm(A @ c_inv - y, np.inf)}")
9 print(f"cond(A) = {cond(A)}")
```

```
error = 6.348273018375039e-10, error using inv(A) = 4.55108965979889e-06
cond(A) = 4462823910804.094
```

## Matrix Inverses Fail Completely for $N = 20$

```
1 N = 20
2 x = np.linspace(0.0, 10.0, N + 1)
3 y = np.exp(x) # example function to interpolate
4 A = np.array([[x_i**n for n in range(N + 1)] for x_i in x]) # or np.vander
5 c = solve(A, y)
6 c_inv = inv(A) @ y # Solving with inv(A) instead of solve(A, y)
7 print(f"error = {norm(A @ c - y, np.inf)}, \
8 error using inv(A) = {norm(A @ c_inv - y, np.inf)}")
9 print(f"cond(A) = {cond(A):.4g}")
```

```
error = 1.9554136088117957e-10, error using inv(A) = 21804.714723170073
cond(A) = 3.325e+24
```

## Moral of this Story

- Use `solve`, which is faster and can often solve ill-conditioned problems. Rarely use `inv`, and only when you know the problem is well-conditioned
- Check conditioning of matrices when doing numerical work as an occasional diagnostic, as it is a good indicator of potential problems and collinearity
- For approximation, never use a monomial basis for polynomials
  - Prefer polynomials like Chebyshev, which are designed to be as orthogonal as possible

```
1 N = 40
2 x = np.linspace(-1, 1, N+1) # Or any other range of x values
3 A = np.array([np.polynomial.Chebyshev.basis(n)(x_i) for n in range(N+1)]
4 A_monomial = np.array([[x_i**n for n in range(N + 1)] for x_i in x]) # or
5 print(f"cond(A) = {cond(A):.4g}, cond(A_monimial) = {cond(A_monomial):.4g}")
```

```
cond(A) = 3.64e+09, cond(A_monimial) = 8.903e+17
```