

ECON526: Quantitative Economics with Data Science Applications

Foundations of Numerical Linear Algebra

Jesse Perla

Table of contents i

Basic Linear Algebra

Solving Linear Systems of Equations

Eigenvalues and Eigenvectors

Least Squares and the Normal Equations

Going Beyond “`reg y x, robust`”

- Data science, econometrics, and macroeconomics are built on linear algebra.
- Numerical linear algebra has all sorts of pitfalls, which become more critical as we scale up to larger problems.
- Speed differences in choosing better algorithms can be orders of magnitude.
- Crucial to know what goes on under-the-hood in Stata/R/python packages for applied work, even if you don't implement it yourself.
- Material here is related to
 - QuantEcon Python
 - QuantEcon Data Science
 - A First Course in Quantitative Economics with Python

Packages

This section uses the following packages:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import scipy
4 from numpy.linalg import cond, matrix_rank, norm
5 from scipy.linalg import inv, solve, det, eig, lu, eigvals
6 from scipy.linalg import solve_triangular, eigvalsh, cholesky
```

Definition (Big-O Notation)

For a function $f(N)$ and a positive constant C , we say $f(N)$ is $O(g(N))$, if there exist positive constants C and N_0 such that:

$$0 \leq f(N) \leq C \cdot g(N) \quad \text{for all } N \geq N_0$$

- Often crucial to know how problems scale asymptotically (as $N \rightarrow \infty$)
- Caution: $f_1(N) = N^3 + N$ is $O(N^3)$ and $f_2(N) = 1000N^2 + 3N$ is $O(N^2)$
 - Asymptotically choose $f_2(N)$ algorithm, but choose $f_1(N)$ for small N
- Simple examples:
 - $x \cdot y = \sum_{n=1}^N x_n y_n$ is $O(N)$ since it requires N multiplications and additions
 - Ax for $A \in \mathbb{R}^{N \times N}$, $x \in \mathbb{R}^N$ is $O(N^2)$ since it requires N dot products, each $O(N)$

Numerical Precision

Computers have finite precision. 64-bit typical, but 32-bit on GPUs

Definition (Machine Epsilon)

For a given datatype, ϵ is defined as $\epsilon = \min_{\delta > 0} \{\delta : 1 + \delta > 1\}$

```
1 print(f"machine epsilon for float64 = {np.finfo(float).eps}")
2 print(f"1 + eps/2 == 1? {1.0 + 1.1e-16 == 1.0}")
3 print(f"machine epsilon for float32 = {np.finfo(np.float32).eps}")
```

```
machine epsilon for float64 = 2.220446049250313e-16
```

```
1 + eps/2 == 1? True
```

```
machine epsilon for float32 = 1.1920928955078125e-07
```

Basic Linear Algebra

Norms

- Common measure of “size” is the Euclidean norm, or L^2 norm for $x \in \mathbb{R}^2$
- Complexity is $O(N)$: requires squaring N times then N additions to sum. Not nested

$$\|x\|_2 = \sqrt{\sum_{n=1}^N x_n^2}$$

```
1 x = np.array([1, 2, 3]) # Calculating different ways (in order of preference)
2 print(np.sqrt(sum(xval**2 for xval in x))) # manual with comprehensions
3 print(np.sqrt(np.sum(np.square(x)))) # broadcasts
4 print(norm(x)) # built-in to numpy norm(x, ord=2) alternatively
5 print(f"||x||_2^2 = {norm(x)**2} = {x.T @ x} = {np.dot(x, x)}")
```

3.7416573867739413

3.7416573867739413

3.7416573867739413

Solving Systems of Equations

- Solving $Ax = b$ for x is equivalent $A^{-1}Ax = A^{-1}b$
- Then since $A^{-1}A = I$, and $Ix = x$, we have $x = A^{-1}b$
- Careful since matrix algebra is not commutative!

```
1 A = np.array([[0, 2], [3, 4]]) # or ((0, 2), (3, 4))
2 b = np.array([2, 1]) # Column vector
3 x = solve(A, b) # Solve Ax = b for x
4 x
```

```
array([-1.,  1.])
```

Using the Inverse Directly

- Can replace the `solve` with a calculation of an inverse
- But it can be slower or less accurate than solving the system directly

```
1 A_inv = inv(A)
2 A_inv @ b # i.e,  $A^{-1} * b$ 
```

```
array([-1.,  1.])
```

Linear Combinations

We can think of solving a system as finding the linear combination of columns of A that equal b

```
1 b_star = x[0] * A[:, 0] + x[1] * A[:, 1] # using x solution
2 print(f"b = {b}, b_star = {b_star}")
```

```
b = [2 1], b_star = [2. 1.]
```

Column Space and Rank

- The column space of a matrix represents all possible linear combinations of its columns.
- It forms a basis for the space of solutions when solving systems of linear equations represented by the matrix
- The rank of a matrix is the dimension of its column space

```
1 A = np.array([[0, 2], [3, 4]])  
2 matrix_rank(A)
```

2

Hence, can solve $Ax = b$ for any $b \in \mathbb{R}^2$ since the column space is the entire space \mathbb{R}^2

Singular Matrices

On the other hand, note

```
1 A = np.array([[1, 2],  
2               [2, 4]])  
3 matrix_rank(A)
```

1

So we can only solve $Ax = b$ for $b \propto \begin{bmatrix} 1 \\ 2 \end{bmatrix} \propto \begin{bmatrix} 2 \\ 4 \end{bmatrix}$

Checking Singularity

```
1 A = np.array([[1, 2], [2, 4]])
2 # An (expensive) way to check if A is singular is if det(A) = 0
3 print(det(A) == 0.0)
4 print(matrix_rank(A) != A.shape[0]) # or check rank
5 # Check before inverting or use exceptions
6 try:
7     inv(A)
8     print("Matrix is not singular (invertible).")
9 except np.linalg.LinAlgError:
10    print("Matrix is singular (non-invertible).")
```

True

True

Matrix is singular (non-invertible).

Determinant is Not Scale Invariant

- Reminder: numerical precision in calculations makes it hard to compare to zero
- The determinate is useful but depends on the scale of the matrix
- A more robust alternative is the condition number (more next lecture)

```
1 eps, K = 1e-8, 100000
2 A = np.array([[1, 2], [1 + eps, 2 + eps]])
3 print(f"det(A)={det(A):.5g}, det(K*A)={det(K*A):.5g}")
4 print(f"cond(A)={cond(A):.5g}, cond(K*A)={cond(A):.5g},")
5 print(f"det(inv(A))={det(inv(A)):.5g}, cond(inv(A))={cond(inv(A)):.5g}")
```

```
det(A)=-1e-08, det(K*A)=-100
cond(A)=1e+09, cond(K*A)=1e+09,
det(inv(A))=-1e+08, cond(inv(A))=1e+09
```

Solving Linear Systems of Equations

Solving Systems with Multiple RHS

- Inverse is nice because you can reuse the A^{-1} to solve $Ax = b$ for many b
- However, you can do this with `solve` as well
- Or can reuse LR factorizations (discussed next)

```
1 A = np.array([[0, 2], [3, 4]])
2 B = np.array([[2,3], [1,2]]) # [2,1] and [3,2] as columns
3 # or: B = np.column_stack([np.array([2, 1]),np.array([3,2])])
4 X = solve(A, B) # Solve AX = B for X
5 print(X)
6 print(f"Checking: A*{X[:,0]} = {A@X[:, 0]} = {B[:,0]}, column of B")
```

```
[[ -1.          -1.33333333]
 [ 1.           1.5         ]]
```

```
Checking: A*[-1.  1.] = [2. 1.] = [2 1], column of B
```

LU(P) Decompositions

- We can “factor” any square A into $PA = LU$ for triangular L and U . Invertible can have $A = LU$, called the LU decomposition. “P” is for partial-pivoting
- Singular matrices may not have full-rank L or U matrices

```
1 A = np.array([[1, 2], [2, 4]])
2 P, L, U = lu(A)
3 print(f"L*U =\n{L @ U}")
4 print(f"P*A =\n{P @ A}")
```

```
L*U =
[[2. 4.]
 [1. 2.]]
```

```
P*A =
[[2. 4.]
 [1. 2.]]
```

P, U, and L

The P matrix is a permutation matrix of “pivots” the others are triangular

```
1 print(f"P =\n{P}")
2 print(f"L =\n{L}")
3 print(f"U =\n{U}")
```

P =

```
[[0.  1.]
 [1.  0.]]
```

L =

```
[[1.  0. ]
 [0.5 1. ]]
```

U =

```
[[2.  4.]
 [0.  0.]]
```

LU Decompositions and Systems of Equations

- Pivoting is typically implied when talking about “LU”
- Used in the default `solve` algorithm (without more structure)
- Solving systems of equations with triangular matrices: for $Ax = LUx = b$
 1. Define $y = Ux$
 2. Solve $Ly = b$ for y and $Ux = y$ for x
- Since both are triangular, process is $O(N^2)$ (but LU itself $O(N^3)$)
- Could be used to find `inv`
 - $A = LU$ then $AA^{-1} = I = LUA^{-1} = I$
 - Solve for Y in $LY = I$, then solve $UA^{-1} = Y$
- Tight connection to textbook Gaussian elimination (including pivoting)

LU for Non-Singular Matrices

```
1 A = np.array([[1, 2], [3, 4]])
2 P, L, U = lu(A)
3 print(f"L*U =\n{L @ U}")
4 print(f"P*A =\n{P @ A}")
```

```
L*U =
[[3. 4.]
 [1. 2.]]
```

```
P*A =
[[3. 4.]
 [1. 2.]]
```

L, U, P

```
1 print(f"P =\n{P}")
2 print(f"L =\n{L}")
3 print(f"U =\n{U}")
```

P =

[[0. 1.]

[1. 0.]]

L =

[[1. 0.]

[0.33333333 1.]]

U =

[[3. 4.]

[0. 0.66666667]]

Backwards Substitution Example

$$Ux = b$$

$$U \equiv \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}, \quad b = \begin{bmatrix} 7 \\ 2 \end{bmatrix}$$

Solving bottom row for x_2

$$2x_2 = 2, \quad x_2 = 1$$

Move up a row, solving for x_1 , substituting for x_2

$$3x_1 + 1x_2 = 7, \quad 3x_1 + 1 \times 1 = 7, \quad x_1 = 2$$

Generalizes to many rows. For L it is “forward substitution”

Use Triangular Structure if Possible

- Triangular matrices of size N can be solved with back substitution in $O(N^2)$
- Is $O(N^2)$ good or bad? Beats, $O(N^3)$ typical of general methods

```
1 U = np.array([[3, 1],  
2               [0, 2]])  
3 b = np.array([7,2])  
4 solve(U,b) # works, but internally does an LU which is  $O(N^3)$   
5 solve_triangular(U, b, lower=False) # fast  $O(N^2)$ 
```

```
array([2., 1.])
```


Symmetric Matrix Structure

Another common matrix type are symmetric, $A = A^T$

```
1 A = np.array([[1, 2], [2, 5]]) # also posdef, not singular
2 b = np.array([1,4])
3 print(f"Symmetric A? {scipy.linalg.issymmetric(A)}")
4 solve(A, b, assume_a="sym") # could also use "pos" since positive definite
```

Symmetric A? True

array([-3., 2.])

Positive Definite Matrices

- A symmetric matrix A is positive definite if $x^T A x > 0$ for all $x \neq 0$
- Useful in many areas, such as covariance matrices. Example

```
1 A = np.array([[1, 2], [2, 5]])  
2 x = np.array([0, 1]) # can't really check for all x  
3 print(f"x^T A x = {x.T @ A @ x}")
```

$x^T A x = 5$

- Example of a symmetric matrix that is not positive definite

```
1 A = np.array([[1, 2], [2, 0]])  
2 print(f"x^T A x = {x.T @ A @ x}") # one counterexample is enough
```

$x^T A x = 0$

- We can check these with eigenvalues

Cholesky Decomposition

- For symmetric positive definite matrices: $L = U^T$,
- Called a Cholesky decomposition: $A = LL^T$ for a lower triangular matrix L .
- Equivalently, could find $A = U^TU$ for an upper triangular matrix U

```
1 A = np.array([[1, 2], [2, 5]])
2 L = cholesky(A, lower=True) # cholesky also defined for upper=True
3 print(L)
4 print(f"L*L^T =\n{L @ L.T}")
```

```
[[1. 0.]
 [2. 1.]]
L*L^T =
[[1. 2.]
 [2. 5.]]
```

Solving Positive Definite Systems

```
1 A = np.array([[1, 2], [2, 5]])
2 b = np.array([1,4])
3 print(solve(A, b, assume_a="pos")) # uses cholesky internally
4
5 L = cholesky(A, lower=True)
6 y = solve_triangular(L, b, lower=True)
7 x = solve_triangular(L.T, y, lower=False)
8 print(x)
```

[-3. 2.]

[-3. 2.]

Cholesky for Covariance Matrices

- Covariance matrices are positive-definite, semi-definite if degenerate
- Key property of Gaussian random variables:
 - $X \sim N(\mu, \Sigma)$ for $\mu \in \mathbb{R}^N, \Sigma \in \mathbb{R}^{N \times N}$
 - $X = \mu + AZ$ for $Z \sim N(0_N, I_N)$ where $AA^T = \Sigma$
- That is, A is the Cholesky decomposition of the covariance matrix

Matrices as Linear Transformations

- Recall: for $x \in \mathbb{R}^N$ we should think of a $f(x) = Ax$ for $A \in \mathbb{R}^{M \times N}$ as a linear transformation from \mathbb{R}^N to \mathbb{R}^M
 - Definition of Linear: $f(ax_1 + bx_2) = af(x_1) + bf(x_2)$ for scalar a, b
- Similarly, the $y = f(x) = Ax$ then $f^{-1}(y) = A^{-1}y$ transforms from \mathbb{R}^M to \mathbb{R}^N
 - If the matrix is square and invertible, this means we can go back and forth without losing information (i.e., it is bijective). Otherwise we may be projected onto a lower-dimensional “manifold” (e.g. a line in a plane or plane in a cube)
- The vector norm $\|x\|_2$ is an important feature in many applications
 - Hence $\|f(x)\|_2 = \|Ax\|_2$ frequently comes up in economics and datascience
 - e.g. linear regression is written as minimizing a vector norm $\|y - X\beta\|_2$
- Matrix structure or decompositions of A help us better understand the $f(x)$ mapping

Definition (Orthogonal Matrices)

A square matrix Q is orthogonal if $Q^{-1} = Q^T$, and hence $Q^T Q = Q Q^T = I$

- For orthogonal Q , $f(x) = Qx$ is interpreted as rotating x without stretching it
- Similarly if $y = f(x) = Qx$ then $f^{-1}(y) = Q^{-1}y = Q^T y$ is rotating y back
- Columns are orthonormal: $Q = \begin{bmatrix} | & & | \\ q_1 & \dots & q_N \\ | & & | \end{bmatrix}$ then
 - $q_i \cdot q_j = 0$ for $i \neq j$ and $q_i \cdot q_i = 1$
- Rotation means the length doesn't change: $\|Qx\|_2 = \|x\|_2$
- Transformations which preserve norms are central in many applications within data science, machine learning, and economics - especially with high-dimensional data

Eigenvalues and Eigenvectors

Eigenvalues and Eigenvectors

- For a square A , an eigenvector x and eigenvalue λ satisfy

$$Ax = \lambda x$$

- $A \in \mathbb{R}^{N \times N}$ has N eigenvalue/eigenvector pairs, possible multiplicity of λ
- Intuition: x is a direction $Ax \propto x$ and λ says how much it “stretches”
- Properties:
 - For any eigenvector x and scalar c then $cx \propto Ax$ as well
 - Symmetric matrices have real eigenvalues and orthogonal eigenvectors.
i.e. $x_1 \cdot x_2 = 0$ for $x_1 \neq x_2$ eigenvectors. Complex in general
 - Singular if and only if it has an eigenvalue of zero
 - Positive (semi)definite if and only if all eigenvalues are strictly (weakly) positive
 - Diagonal matrix has eigenvalues as its diagonal
 - Triangular matrix has eigenvalues as its diagonal

Positive Definite and Eigenvalues

You cannot check $x^T Ax > 0$ for all x . Check if “stretching” is positive

```
1 A = np.array([[3, 1], [2, 1]])
2 # A_eigs = np.real(eigvals(A)) # symmetric matrices have real eigenvalues
3 A_eigs = eigvalsh(A) # specialized for symmetric/hermitian matrices
4 print(A_eigs)
5 is_positive_definite = np.all(A_eigs > 0)
6 is_positive_semi_definite = np.all(A_eigs >= 0) # or eigvals(A) >= -eps
7 print(f"pos-def? {is_positive_definite}")
8 print(f"pos-semi-def? {is_positive_semi_definite}")
```

```
[-0.23606798  4.23606798]
```

```
pos-def? False
```

```
pos-semi-def? False
```

Positive Semi-Definite Matrices May Have a Zero Eigenvalue

The simplest positive-semi-definite (but not posdef) matrix is

```
1 A_eigs = eigvalsh(np.array([[1, 0], [0, 0]]))
2 print(A_eigs)
3 is_positive_definite = np.all(A_eigs > 0)
4 is_positive_semi_definite = np.all(A_eigs >= 0) # or eigvals(A) >= -eps
5 print(f"pos-def? {is_positive_definite}")
6 print(f"pos-semi-def? {is_positive_semi_definite}")
```

[0. 1.]

pos-def? False

pos-semi-def? True

Eigen Decomposition (i.e., Spectral Decomposition if Symmetric)

- For square, symmetric, non-singular matrix A factor into

$$A = Q\Lambda Q^{-1}$$

- Q is a matrix of eigenvectors, Λ is a diagonal matrix of paired eigenvalues (in order)
- For symmetric matrices, the eigenvectors are orthogonal and $Q^{-1}Q = Q^TQ = I$, and we can think of them as forming an orthonormal basis
- Orthogonal matrices can be thought of as rotations without stretching
- More general matrices all have a Singular Value Decomposition (SVD)
- With symmetric A , an interpretation of Ax is that we can first rotate x into the Q basis, then stretch by Λ , then rotate back
- Can be used to find A^t for large t (e.g. for Markov chains)
 - P^t , i.e. $P \cdot P \cdot \dots \cdot P$ for t times

Spectral/Eigendecomposition of Symmetric Matrix Example

```
1 A = np.array([[2, 1], [1, 3]])
2 Lambda, Q = eig(A)
3 print(f"eigenvectors are column-by-column in Q =\n{Q}")
4 print(f"eigenvalues are in Lambda = {Lambda}")
5 print(f"Q Lambda Q^T =\n{Q @ np.diag(np.real(Lambda)) @ Q.T}")
```

eigenvectors are column-by-column in Q =

```
[[ -0.85065081 -0.52573111]
```

```
 [ 0.52573111 -0.85065081]]
```

eigenvalues are in Lambda = [1.38196601+0.j 3.61803399+0.j]

Q Lambda Q^T =

```
[[2. 1.]
```

```
 [1. 3.]]
```

Spectral Radius is Maximum Absolute Eigenvalue

- If any $\lambda \in \Lambda$ are > 1 can see this would explode
- Useful for seeing if iteration $x_{t+1} = Ax_t$ from a x_0 explodes

Definition (Spectral Radius)

The spectral radius of matrix A is $\rho(A) = \max_{\lambda \in \Lambda} |\lambda|$

Least Squares and the Normal Equations

Least Squares

Given a matrix $X \in \mathbb{R}^{N \times M}$ and a vector $y \in \mathbb{R}^N$, we want to find $\beta \in \mathbb{R}^M$ such that

$$\min_{\beta} \|y - X\beta\|^2, \text{ that is,}$$

$$\min_{\beta} \sum_{n=1}^N \frac{1}{N} (y_n - X_n \cdot \beta)^2$$

Where X_n is n'th row. Take FOCS and rearrange to get

$$(X^T X)\beta = X^T y$$

Solving the Normal Equations

- The X is often referred to as the “design matrix”. $X^T X$ as the Gram matrix
- Can form $A = X^T X$ and $b = X^T y$ and solve $A\beta = b$.
 - Or invert $X^T X$ to get $\beta = (X^T X)^{-1} X^T y$
 - Note that $X^T X$ is symmetric and, if X is full-rank, positive definite
- In practice, use the `lstsq` function in `scipy`
 - It uses better algorithms using eigenvectors. More stable (see next lecture on conditioning)
 - One algorithm uses another factoring, the QR decomposition
 - There, $X = QR$ for Q orthogonal and R upper triangular. See QR Decomposition for more
- We are exploring linear algebra in this lecture
 - For applied work use higher-level libraries like `statsmodels` (integrated well with `pandas` and `seaborn`)
 - See `statsmodels` docs for R-style notation
 - See QuantEcon OLS Notes for more.

Example of LLS using Scipy

```
1 N, M = 100, 5
2 X = np.random.randn(N, M)
3 beta = np.random.randn(M)
4 y = X @ beta + 0.05 * np.random.randn(N)
5 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X, y)
6 print(f"beta =\n {beta}\nbeta_hat =\n{beta_hat}")
```

beta =

[-1.56627442 -0.38707104 -0.53992938 1.88893498 -0.12611984]

beta_hat =

[-1.57370971 -0.3923169 -0.53467955 1.88647361 -0.12421782]

Solving using the Normal Equations

Or we can solve it directly. Provide matrix structure (so it can use a Cholesky)

```
1 beta_hat = solve(X.T @ X, X.T @ y, assume_a="pos")
2 print(f"beta =\n {beta}\nbeta_hat =\n{beta_hat}")
```

beta =

`[-1.56627442 -0.38707104 -0.53992938 1.88893498 -0.12611984]`

beta_hat =

`[-1.57370971 -0.3923169 -0.53467955 1.88647361 -0.12421782]`

Collinearity in “Tall” Matrices

- Tall $\mathbb{R}^{N \times M}$ “design matrices” have $N > M$ and are “overdetermined”
- The rank of a matrix is full rank if all columns are linearly independent
- You can only identify M parameters with M linearly independent columns

```
1 X = np.array([[1, 2], [2, 5], [3, 7]]) # 3 observations, 2 variables
2 X_col = np.array([[1, 2], [2, 4], [3, 6]]) # all proportional
3 print(f"rank(X) = {matrix_rank(X)}, rank(X_col) = {matrix_rank(X_col)}")
```

rank(X) = 2, rank(X_col) = 1

Collinearity and Estimation

- If X is not full rank, then $X^T X$ is not invertible. For example:

```
1 print(f"cond(X'*X)={cond(X.T@X)}, cond(X_col'*X_col)={cond(X_col.T@X_col)}")
```

```
cond(X'*X)=2819.3329786399063, cond(X_col'*X_col)=1.2999933999712892e+16
```

- Note that when you start doing operations on matrices, numerical error creeps in, so you will not get an exact number
- The rule-of-thumb with condition numbers is that if it is 1×10^k then you lose about k digits of precision. So this effectively means it is singular
- Given the singular matrix, this means a continuum of β will solve the problem

lstsq Solves it? Careful on Interpretation!

- Since $X_{col}^T X_{col}$ is singular, we cannot use `solve(X.T@X, y)`
- But what about `lstsq` methods?
- As you will see, this gives an answer. Interpretation is hard
- The key is that in the case of non-full rank, you cannot identify individual parameters
 - Related to “Identification” in econometrics
 - Having low residuals is not enough

```
1 y = np.array([5.0, 10.1, 14.9])
2 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X_col, y)
3 print(f"beta_hat_col = {beta_hat}")
4 print(f"rank={rank}, cols={X.shape[1]}, norm(X*beta_hat_col-y)={norm(residuals)}
```

```
beta_hat_col = [0.99857143  1.99714286]
rank=1, cols=2, norm(X*beta_hat_col-y)=0.0
```

Fat Design Matrices

- Fat $\mathbb{R}^{N \times M}$ “design matrices” have $N < M$ and are “underdetermined”
- Less common in econometrics, but useful to understand the structure
- A continuum $\beta \in \mathbb{R}^{M - \text{rank}(X)}$ solve this problem

```
1 X = np.array([[1, 2, 3], [0, 5, 7]]) # 2 rows, 3 variables
2 y = np.array([5, 10])
3 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X, y)
4 print(f"beta_hat = {beta_hat}, rank={rank}, ? residuals = {residuals}")
```

```
beta_hat = [0.8 0.6 1. ], rank=2, ? residuals = []
```

Which Solution?

- Residuals are zero here because there are enough parameters to fit perfectly (i.e., it is underdetermined)
- Given the multiple solutions, the `lstsq` is giving

$$\min_{\beta} \|\beta\|_2^2 \text{ s.t. } X\beta = y$$

- i.e., the “smallest” coefficients which interpolate the data exactly
- Which trivially fulfills the OLS objective: $\min_{\beta} \|y - X\beta\|_2^2$
- Useful and common in ML, but be **very** careful when interpreting for economics
 - Tight connections to Bayesian versions of statistical tests
 - But until you understand econometrics and “identification” well, **stick to full-rank matrices**
 - **Advanced topics:** search for “Regularization”, “Ridgeless Regression” and “Benign Overfitting in Linear Regression.”