

ECON526: Quantitative Economics with Data Science Applications

Foundations of Numerical Linear Algebra

Jesse Perla

Basic Linear Algebra

Solving Linear Systems of Equations

Eigenvalues and Eigenvectors

Least Squares and the Normal Equations

Motivation and Materials

- Data science, econometrics, and macroeconomics are built on linear algebra.
- Numerical linear algebra has all sorts of pitfalls, which become more critical as we scale up to larger problems.
- Speed differences in choosing better algorithms can be orders of magnitude.
- Crucial to know what goes on under-the-hood in Stata/R/python packages for applied work, even if you don't implement it yourself.
- Material here is related to
 - QuantEcon Python
 - QuantEcon Data Science
 - A First Course in Quantitative Economics with Python

This section uses the following packages:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import scipy
4 from numpy.linalg import cond, matrix_rank, norm
5 from scipy.linalg import inv, solve, det, eig, lu, eigvals
6 from scipy.linalg import solve_triangular, eigvalsh, cholesky
```

Definition (Big-O Notation)

For a function $f(N)$ and a positive constant C , we say $f(N)$ is $O(g(N))$, if there exist positive constants C and N_0 such that:

$$0 \leq f(N) \leq C \cdot g(N) \quad \text{for all } N \geq N_0$$

- Often crucial to know how problems scale asymptotically (as $N \rightarrow \infty$)
- Caution: $f_1(N) = N^3 + N$ is $O(N^3)$ and $f_2(N) = 1000N^2 + 3N$ is $O(N^2)$
 - Asymptotically choose $f_2(N)$ algorithm, but choose $f_1(N)$ for small N
- Simple examples:
 - $x \cdot y = \sum_{n=1}^N x_n y_n$ is $O(N)$ since it requires N multiplications and additions
 - Ax for $A \in \mathbb{R}^{N \times N}$, $x \in \mathbb{R}^N$ is $O(N^2)$ since it requires N dot products, each $O(N)$

Numerical Precision

Computers have finite precision. 64-bit typical, but 32-bit on GPUs

Definition (Machine Epsilon)

For a given datatype, ϵ is defined as $\epsilon = \min_{\delta > 0} \{\delta : 1 + \delta > 1\}$

```
1 print(f"machine epsilon for float64 = {np.finfo(float).eps}")
2 print(f"1 + eps/2 == 1? {1.0 + 1.1e-16 == 1.0}")
3 print(f"machine epsilon for float32 = {np.finfo(np.float32).eps}")
```

```
machine epsilon for float64 = 2.220446049250313e-16
```

```
1 + eps/2 == 1? True
```

```
machine epsilon for float32 = 1.1920928955078125e-07
```

Basic Linear Algebra

Norms

- Common measure of “size” is the Euclidean norm, or L^2 norm for $x \in \mathbb{R}^2$
- Complexity is $O(N)$: requires squaring N times then N additions to sum. Not nested

$$\|x\|_2 = \sqrt{\sum_{n=1}^N x_n^2}$$

```
1 x = np.array([1, 2, 3]) # Calculating different ways (in order of preference)
2 print(np.sqrt(sum(xval**2 for xval in x))) # manual with comprehensions
3 print(np.sqrt(np.sum(np.square(x)))) # broadcasts
4 print(norm(x)) # built-in to numpy norm(x, ord=2) alternatively
5 print(f"||x||_2^2 = {norm(x)**2} = {x.T @ x} = {np.dot(x, x)}")
```

3.7416573867739413

3.7416573867739413

3.7416573867739413

||x||_2^2 = 14.0 = 14 = 14

Solving Systems of Equations

- Solving $Ax = b$ for x is equivalent $A^{-1}Ax = A^{-1}b$
- Then since $A^{-1}A = I$, and $Ix = x$, we have $x = A^{-1}b$
- Careful since matrix algebra is not commutative!

```
1 A = np.array([[0, 2], [3, 4]]) # or ((0, 2), (3, 4))
2 b = np.array([2, 1]) # Column vector
3 x = solve(A, b) # Solve Ax = b for x
4 x
```

```
array([-1.,  1.])
```

Using the Inverse Directly

- Can replace the `solve` with a calculation of an inverse
- But it can be slower or less accurate than solving the system directly

```
1 A_inv = inv(A)
2 A_inv @ b # i.e,  $A^{-1} * b$ 
```

```
array([-1.,  1.])
```

We can think of solving a system as finding the linear combination of columns of A that equal b

```
1 b_star = x[0] * A[:, 0] + x[1] * A[:, 1] # using x solution
2 print(f"b = {b}, b_star = {b_star}")
```

```
b = [2 1], b_star = [2. 1.]
```

Column Space and Rank

- The column space of a matrix represents all possible linear combinations of its columns.
- It forms a basis for the space of solutions when solving systems of linear equations represented by the matrix
- The rank of a matrix is the dimension of its column space

```
1 A = np.array([[0, 2], [3, 4]])  
2 matrix_rank(A)
```

2

Hence, can solve $Ax = b$ for any $b \in \mathbb{R}^2$ since the column space is the entire space \mathbb{R}^2

On the other hand, note

```
1 A = np.array([[1, 2],  
2               [2, 4]])  
3 matrix_rank(A)
```

1

So we can only solve $Ax = b$ for $b \propto \begin{bmatrix} 1 \\ 2 \end{bmatrix} \propto \begin{bmatrix} 2 \\ 4 \end{bmatrix}$

Checking Singularity

```
1 A = np.array([[1, 2], [2, 4]])
2 # An (expensive) way to check if A is singular is if det(A) = 0
3 print(det(A) == 0.0)
4 print(matrix_rank(A) != A.shape[0]) # or check rank
5 # Check before inverting or use exceptions
6 try:
7     inv(A)
8     print("Matrix is not singular (invertible).")
9 except np.linalg.LinAlgError:
10    print("Matrix is singular (non-invertible).")
```

True

True

Matrix is singular (non-invertible).

Determinant is Not Scale Invariant

- Reminder: numerical precision in calculations makes it hard to compare to zero
- The determinate is useful but depends on the scale of the matrix
- A more robust alternative is the condition number (more next lecture)

```
1 eps, K = 1e-8, 100000
2 A = np.array([[1, 2], [1 + eps, 2 + eps]])
3 print(f"det(A)={det(A):.5g}, det(K*A)={det(K*A):.5g}")
4 print(f"cond(A)={cond(A):.5g}, cond(K*A)={cond(A):.5g},")
5 print(f"det(inv(A))={det(inv(A)):.5g}, cond(inv(A))={cond(inv(A)):.5g}")
```

```
det(A)=-1e-08, det(K*A)=-100
cond(A)=1e+09, cond(K*A)=1e+09,
det(inv(A))=-1e+08, cond(inv(A))=1e+09
```

Solving Linear Systems of Equations

Solving Systems with Multiple RHS

- Inverse is nice because you can reuse the A^{-1} to solve $Ax = b$ for many b
- However, you can do this with `solve` as well
- Or can reuse LR factorizations (discussed next)

```
1 A = np.array([[0, 2], [3, 4]])
2 B = np.array([[2,3], [1,2]]) # [2,1] and [3,2] as columns
3 # or: B = np.column_stack([np.array([2, 1]),np.array([3,2])])
4 X = solve(A, B) # Solve AX = B for X
5 print(X)
6 print(f"Checking: A*{X[:,0]} = {A@X[:, 0]} = {B[:,0]}, column of B")
```

```
[[ -1.          -1.33333333]
 [  1.           1.5         ]]
```

```
Checking: A*[-1.  1.] = [2. 1.] = [2 1], column of B
```

LU(P) Decompositions

- We can “factor” any square A into $PA = LU$ for triangular L and U . Invertible can have $A = LU$, called the LU decomposition. “P” is for partial-pivoting
- Singular matrices may not have full-rank L or U matrices

```
1 A = np.array([[1, 2], [2, 4]])  
2 P, L, U = lu(A)  
3 print(f"L*U =\n{L @ U}")  
4 print(f"P*A =\n{P @ A}")
```

L*U =

```
[[2. 4.]  
 [1. 2.]]
```

P*A =

```
[[2. 4.]  
 [1. 2.]]
```

The P matrix is a permutation matrix of “pivots” the others are triangular

```
1 print(f"P =\n{P}")
2 print(f"L =\n{L}")
3 print(f"U =\n{U}")
```

P =

```
[[0. 1.]
 [1. 0.]]
```

L =

```
[[1.  0. ]
 [0.5 1. ]]
```

U =

```
[[2. 4.]
 [0. 0.]]
```

LU Decompositions and Systems of Equations

- Pivoting is typically implied when talking about “LU”
- Used in the default `solve` algorithm (without more structure)
- Solving systems of equations with triangular matrices: for $Ax = LUx = b$
 1. Define $y = Ux$
 2. Solve $Ly = b$ for y and $Ux = y$ for x
- Since both are triangular, process is $O(N^2)$ (but LU itself $O(N^3)$)
- Could be used to find `inv`
 - $A = LU$ then $AA^{-1} = I = LUA^{-1} = I$
 - Solve for Y in $LY = I$, then solve $UA^{-1} = Y$
- Tight connection to textbook Gaussian elimination (including pivoting)

LU for Non-Singular Matrices

```
1 A = np.array([[1, 2], [3, 4]])
2 P, L, U = lu(A)
3 print(f"L*U =\n{L @ U}")
4 print(f"P*A =\n{P @ A}")
```

L*U =

```
[[3. 4.]
 [1. 2.]]
```

P*A =

```
[[3. 4.]
 [1. 2.]]
```

L, U, P

```
1 print(f"P =\n{P}")
2 print(f"L =\n{L}")
3 print(f"U =\n{U}")
```

P =

```
[[0. 1.]
 [1. 0.]]
```

L =

```
[[1.          0.          ]
 [0.33333333  1.          ]]
```

U =

```
[[3.          4.          ]
 [0.          0.66666667]]
```

Backwards Substitution Example

$$Ux = b$$

$$U \equiv \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}, \quad b = \begin{bmatrix} 7 \\ 2 \end{bmatrix}$$

Solving bottom row for x_2

$$2x_2 = 2, \quad x_2 = 1$$

Move up a row, solving for x_1 , substituting for x_2

$$3x_1 + 1x_2 = 7, \quad 3x_1 + 1 \times 1 = 7, \quad x_1 = 2$$

Generalizes to many rows. For L it is “forward substitution”

Use Triangular Structure if Possible

- Triangular matrices of size N can be solved with back substitution in $O(N^2)$
- Is $O(N^2)$ good or bad? Beats, $O(N^3)$ typical of general methods

```
1 U = np.array([[3, 1],  
2               [0, 2]])  
3 b = np.array([7, 2])  
4 solve(U, b) # works, but internally does an LU which is  $O(N^3)$   
5 solve_triangular(U, b, lower=False) # fast  $O(N^2)$ 
```

```
array([2., 1.])
```


Symmetric Matrix Structure

Another common matrix type are symmetric, $A = A^T$

```
1 A = np.array([[1, 2], [2, 5]]) # also posdef, not singular
2 b = np.array([1, 4])
3 print(f"Symmetric A? {scipy.linalg.issymmetric(A)}")
4 solve(A, b, assume_a="sym") # could also use "pos" since positive definite
```

Symmetric A? True

array([-3., 2.])

Positive Definite Matrices

- A symmetric matrix A is positive definite if $x^T A x > 0$ for all $x \neq 0$
- Useful in many areas, such as covariance matrices. Example

```
1 A = np.array([[1, 2], [2, 5]])
2 x = np.array([0, 1]) # can't really check for all x
3 print(f"x^T A x = {x.T @ A @ x}")
```

$$x^T A x = 5$$

- Example of a symmetric matrix that is not positive definite

```
1 A = np.array([[1, 2], [2, 0]])
2 print(f"x^T A x = {x.T @ A @ x}") # one counterexample is enough
```

$$x^T A x = 0$$

- We can check these with eigenvalues

Cholesky Decomposition

- For symmetric positive definite matrices: $L = U^T$,
- Called a Cholesky decomposition: $A = LL^T$ for a lower triangular matrix L .
- Equivalently, could find $A = U^T U$ for an upper triangular matrix U

```
1 A = np.array([[1, 2], [2, 5]])
2 L = cholesky(A, lower=True) # cholesky also defined for upper=True
3 print(L)
4 print(f"L*L^T =\n{L @ L.T}")
```

```
[[1. 0.]
 [2. 1.]]
L*L^T =
[[1. 2.]
 [2. 5.]]
```

Solving Positive Definite Systems

```
1 A = np.array([[1, 2], [2, 5]])
2 b = np.array([1, 4])
3 print(solve(A, b, assume_a="pos")) # uses cholesky internally
4
5 L = cholesky(A, lower=True)
6 y = solve_triangular(L, b, lower=True)
7 x = solve_triangular(L.T, y, lower=False)
8 print(x)
```

[-3. 2.]

[-3. 2.]

- Covariance matrices are positive-definite, semi-definite if degenerate
- Key property of Gaussian random variables:
 - $X \sim N(\mu, \Sigma)$ for $\mu \in \mathbb{R}^N, \Sigma \in \mathbb{R}^{N \times N}$
 - $X = \mu + AZ$ for $Z \sim N(0_N, I_N)$ where $AA^T = \Sigma$
- That is, A is the Cholesky decomposition of the covariance matrix

Eigenvalues and Eigenvectors

Eigenvalues and Eigenvectors

- For a square A , an eigenvector x and eigenvalue λ satisfy

$$Ax = \lambda x$$

- $A \in \mathbb{R}^{N \times N}$ has N eigenvalue/eigenvector pairs, possible multiplicity of λ
- Intuition: x is a direction $Ax \propto x$ and λ says how much it “stretches”
- Properties:
 - For any eigenvector x and scalar c then $cx \propto Ax$ as well
 - Symmetric matrices have real eigenvalues and orthogonal eigenvectors. i.e. $x_1 \cdot x_2 = 0$ for $x_1 \neq x_2$ eigenvectors. Complex in general
 - Singular if and only if it has an eigenvalue of zero
 - Positive (semi)definite if and only if all eigenvalues are strictly (weakly) positive
 - Diagonal matrix has eigenvalues as its diagonal
 - Triangular matrix has eigenvalues as its diagonal

Positive Definite and Eigenvalues

You cannot check $x^T Ax > 0$ for all x . Check if “stretching” is positive

```
1 A = np.array([[3, 1], [2, 1]])
2 # A_eigs = np.real(eigvals(A)) # symmetric matrices have real eigenvalues
3 A_eigs = eigvalsh(A) # specialized for symmetric/hermitian matrices
4 print(A_eigs)
5 is_positive_definite = np.all(A_eigs > 0)
6 is_positive_semi_definite = np.all(A_eigs >= 0) # or eigvals(A) >= -eps
7 print(f"pos-def? {is_positive_definite}")
8 print(f"pos-semi-def? {is_positive_semi_definite}")
```

```
[-0.23606798  4.23606798]
```

```
pos-def? False
```

```
pos-semi-def? False
```


Positive Semi-Definite Matrices May Have a Zero Eigenvalue

The simplest positive-semi-definite (but not posdef) matrix is

```
1 A_eigs = eigvalsh(np.array([[1, 0], [0, 0]]))
2 print(A_eigs)
3 is_positive_definite = np.all(A_eigs > 0)
4 is_positive_semi_definite = np.all(A_eigs >= 0) # or eigvals(A) >= -eps
5 print(f"pos-def? {is_positive_definite}")
6 print(f"pos-semi-def? {is_positive_semi_definite}")
```

[0. 1.]

pos-def? False

pos-semi-def? True

Eigenvalue Decomposition

- For square, symmetric, non-singular matrix A factor into

$$A = Q\Lambda Q^{-1}$$

- Q is a matrix of eigenvectors, Λ is a diagonal matrix of eigenvalues (in order)
- For symmetric matrices, the eigenvectors are orthogonal and $Q^{-1}Q = Q^T Q = I$, and we can think of them as forming an orthonormal basis
- Orthogonal matrices can be thought of as rotations without stretching
- More general matrices all have a Singular Value Decomposition (SVD)
- With symmetric A , an interpretation of Ax is that we can first rotate x into the Q basis, then stretch by Λ , then rotate back
- Can be used to find A^t for large t (e.g. for Markov chains)
 - P^t , i.e. $P \cdot P \cdot \dots \cdot P$ for t times
 - $P = Q\Lambda Q^{-1}$ then $P^t = Q\Lambda^t Q^{-1}$ where Λ^t is just the pointwise power

Eigenvalue Decomposition of Symmetric Matrix Example

```
1 A = np.array([[2, 1], [1, 3]])
2 Lambda, Q = eig(A)
3 print(f"eigenvectors are column-by-column in Q =\n{Q}")
4 print(f"eigenvalues are in Lambda = {Lambda}")
5 print(f"Q Lambda Q^T =\n{Q @ np.diag(np.real(Lambda)) @ Q.T}")
```

eigenvectors are column-by-column in Q =

```
[[-0.85065081 -0.52573111]
 [ 0.52573111 -0.85065081]]
```

eigenvalues are in Lambda = [1.38196601+0.j 3.61803399+0.j]

Q Lambda Q^T =

```
[[2. 1.]
 [1. 3.]]
```

Spectral Radius is Maximum Absolute Eigenvalue

- If any $\lambda \in \Lambda$ are > 1 can see this would explode
- Useful for seeing if iteration $x_{t+1} = Ax_t$ from a x_0 explodes

Definition (Spectral Radius)

The spectral radius of matrix A is $\rho(A) = \max_{\lambda \in \Lambda} |\lambda|$

Least Squares and the Normal Equations

Given a matrix $X \in \mathbb{R}^{N \times M}$ and a vector $y \in \mathbb{R}^N$, we want to find $\beta \in \mathbb{R}^M$ such that

$$\min_{\beta} \|y - X\beta\|^2, \text{ that is,}$$
$$\min_{\beta} \sum_{n=1}^N \frac{1}{N} (y_n - X_n \cdot \beta)^2$$

Where X_n is n'th row. Take FOCS and rearrange to get

$$(X^T X)\beta = X^T y$$

Solving the Normal Equations

- The X is often referred to as the “design matrix”. $X^T X$ as the Gram matrix
- Can form $A = X^T X$ and $b = X^T y$ and solve $A\beta = b$.
 - Or invert $X^T X$ to get $\beta = (X^T X)^{-1} X^T y$
 - Note that $X^T X$ is symmetric and, if X is full-rank, positive definite
- In practice, use the `lstsq` function in `scipy`
 - It uses better algorithms using eigenvectors. More stable (see next lecture on conditioning)
 - One algorithm uses another factoring, the QR decomposition
 - There, $X = QR$ for Q orthogonal and R upper triangular. See QR Decomposition for more
- We are exploring linear algebra in this lecture
 - For applied work use higher-level libraries like `statsmodels` (integrated well with `pandas` and `seaborn`)
 - See `statsmodels` docs for R-style notation
 - See QuantEcon OLS for later

Example of LLS using Scipy

```
1 N, M = 100, 5
2 X = np.random.randn(N, M)
3 beta = np.random.randn(M)
4 y = X @ beta + 0.05 * np.random.randn(N)
5 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X, y)
6 print(f"beta =\n {beta}\nbeta_hat =\n{beta_hat}")
```

beta =

```
[ 0.5356111  0.00589317 -0.69347853 -1.32956515 -0.95302447]
```

beta_hat =

```
[ 0.54533529  0.00168887 -0.69943815 -1.33164324 -0.95063408]
```


Solving using the Normal Equations

Or we can solve it directly. Provide matrix structure (so it can use a Cholesky)

```
1 beta_hat = solve(X.T @ X, X.T @ y, assume_a="pos")
2 print(f"beta =\n {beta}\nbeta_hat =\n{beta_hat}")
```

beta =

```
[ 0.5356111  0.00589317 -0.69347853 -1.32956515 -0.95302447]
```

beta_hat =

```
[ 0.54533529  0.00168887 -0.69943815 -1.33164324 -0.95063408]
```

Collinearity in “Tall” Matrices

- Tall $\mathbb{R}^{N \times M}$ “design matrices” have $N > M$ and are “overdetermined”
- The rank of a matrix is full rank if all columns are linearly independent
- You can only identify M parameters with M linearly independent columns

```
1 X = np.array([[1, 2], [2, 5], [3, 7]]) # 3 observations, 2 variables
2 X_col = np.array([[1, 2], [2, 4], [3, 6]]) # all proportional
3 print(f"rank(X) = {matrix_rank(X)}, rank(X_col) = {matrix_rank(X_col)}")
```

`rank(X) = 2, rank(X_col) = 1`

- If X is not full rank, then $X^T X$ is not invertible. For example:

```
1 print(f"cond(X'*X)={cond(X.T@X)}, cond(X_col'*X_col)={cond(X_col.T@X_col)}")
```

```
cond(X'*X)=2819.3329786399063, cond(X_col'*X_col)=1.2999933999712892e+16
```

- Note that when you start doing operations on matrices, numerical error creeps in, so you will not get an exact number
- The rule-of-thumb with condition numbers is that if it is 1×10^k then you lose about k digits of precision. So this effectively means it is singular
- Given the singular matrix, this means a continuum of β will solve the problem

lstsq Solves it? Careful on Interpretation!

- Since $X_{col}^T X_{col}$ is singular, we cannot use `solve(X.T@X, y)`
- But what about `lstsq` methods?
- As you will see, this gives an answer. Interpretation is hard
- The key is that in the case of non-full rank, you cannot identify individual parameters
 - Related to “Identification” in econometrics
 - Having low residuals is not enough

```
1 y = np.array([5.0, 10.1, 14.9])
2 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X_col, y)
3 print(f"beta_hat_col = {beta_hat}")
4 print(f"rank={rank}, cols={X.shape[1]}, norm(X*beta_hat_col-y)={norm(residuals)}
```

```
beta_hat_col = [0.99857143 1.99714286]
rank=1, cols=2, norm(X*beta_hat_col-y)=0.0
```

- Fat $\mathbb{R}^{N \times M}$ “design matrices” have $N < M$ and are “underdetermined”
- Less common in econometrics, but useful to understand the structure
- A continuum $\beta \in \mathbb{R}^{M - \text{rank}(X)}$ solve this problem

```
1 X = np.array([[1, 2, 3], [0, 5, 7]]) # 2 rows, 3 variables
2 y = np.array([5, 10])
3 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X, y)
4 print(f"beta_hat = {beta_hat}, rank={rank}, ? residuals = {residuals}")
```

```
beta_hat = [0.8 0.6 1. ], rank=2, ? residuals = []
```

Which Solution?

- Residuals are zero here because there are enough parameters to fit perfectly (i.e., it is underdetermined)
- Given the multiple solutions, the `lstsq` is giving

$$\min_{\beta} \|\beta\|_2^2 \text{ s.t. } X\beta = y$$

- i.e., the “smallest” coefficients which interpolate the data exactly
- Which trivially fulfills the OLS objective: $\min_{\beta} \|y - X\beta\|_2^2$
- Useful and common in ML, but be **very** careful when interpreting for economics
 - Tight connections to Bayesian versions of statistical tests
 - But until you understand econometrics and “identification” well, **stick to full-rank matrices**
 - **Advanced topics:** search for “Regularization”, “Ridgeless Regression” and “Benign Overfitting in Linear Regression.”