

# ECON526: Quantitative Economics with Data Science Applications

Foundations of Numerical Linear Algebra

---

Jesse Perla

Basic Linear Algebra

Solving Linear Systems of Equations

Eigenvalues and Eigenvectors

Least Squares and the Normal Equations

# Motivation and Materials

- Data science, econometrics, and macroeconomics are built on linear algebra.
- Numerical linear algebra has all sorts of pitfalls, which become more critical as we scale up to larger problems.
- Speed differences in choosing better algorithms can be orders of magnitude.
- Crucial to know what goes on under-the-hood in Stata/R/python packages for applied work, even if you don't implement it yourself.
- Material here is related to
  - QuantEcon Python
  - QuantEcon Data Science
  - A First Course in Quantitative Economics with Python

This section uses the following packages:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import scipy
4 from numpy.linalg import cond, matrix_rank, norm
5 from scipy.linalg import inv, solve, det, eig, lu, eigvals
6 from scipy.linalg import solve_triangular, eigvalsh, cholesky
```

## Definition (Big-O Notation)

For a function  $f(N)$  and a positive constant  $C$ , we say  $f(N)$  is  $O(g(N))$ , if there exist positive constants  $C$  and  $N_0$  such that:

$$0 \leq f(N) \leq C \cdot g(N) \quad \text{for all } N \geq N_0$$

- Often crucial to know how problems scale asymptotically (as  $N \rightarrow \infty$ )
- Caution:  $f_1(N) = N^3 + N$  is  $O(N^3)$  and  $f_2(N) = 1000N^2 + 3N$  is  $O(N^2)$ 
  - Asymptotically choose  $f_2(N)$  algorithm, but choose  $f_1(N)$  for small  $N$
- Simple examples:
  - $x \cdot y = \sum_{n=1}^N x_n y_n$  is  $O(N)$  since it requires  $N$  multiplications and additions
  - $Ax$  for  $A \in \mathbb{R}^{N \times N}$ ,  $x \in \mathbb{R}^N$  is  $O(N^2)$  since it requires  $N$  dot products, each  $O(N)$

# Numerical Precision

Computers have finite precision. 64-bit typical, but 32-bit on GPUs

## Definition (Machine Epsilon)

For a given datatype,  $\epsilon$  is defined as  $\epsilon = \min_{\delta > 0} \{\delta : 1 + \delta > 1\}$

```
1 print(f"machine epsilon for float64 = {np.finfo(float).eps}")
2 print(f"1 + eps/2 == 1? {1.0 + 1.1e-16 == 1.0}")
3 print(f"machine epsilon for float32 = {np.finfo(np.float32).eps}")
```

```
machine epsilon for float64 = 2.220446049250313e-16
```

```
1 + eps/2 == 1? True
```

```
machine epsilon for float32 = 1.1920928955078125e-07
```

# Basic Linear Algebra

---

## Norms

- Common measure of “size” is the Euclidean norm, or  $L^2$  norm for  $x \in \mathbb{R}^2$
- Complexity is  $O(N)$ : requires squaring  $N$  times then  $N$  additions to sum. Not nested

$$\|x\|_2 = \sqrt{\sum_{n=1}^N x_n^2}$$

```
1 x = np.array([1, 2, 3]) # Calculating different ways (in order of preference)
2 print(np.sqrt(sum(xval**2 for xval in x))) # manual with comprehensions
3 print(np.sqrt(np.sum(np.square(x)))) # broadcasts
4 print(norm(x)) # built-in to numpy norm(x, ord=2) alternatively
5 print(f"||x||_2^2 = {norm(x)**2} = {x.T @ x} = {np.dot(x, x)}")
```

3.7416573867739413

3.7416573867739413

3.7416573867739413

||x||\_2^2 = 14.0 = 14 = 14



## Solving Systems of Equations

- Solving  $Ax = b$  for  $x$  is equivalent  $A^{-1}Ax = A^{-1}b$
- Then since  $A^{-1}A = I$ , and  $Ix = x$ , we have  $x = A^{-1}b$
- Careful since matrix algebra is not commutative!

```
1 A = np.array([[0, 2], [3, 4]]) # or ((0, 2), (3, 4))
2 b = np.array([2, 1]) # Column vector
3 x = solve(A, b) # Solve Ax = b for x
4 x
```

```
array([-1.,  1.])
```

## Using the Inverse Directly

- Can replace the `solve` with a calculation of an inverse
- But it can be slower or less accurate than solving the system directly

```
1 A_inv = inv(A)
2 A_inv @ b # i.e,  $A^{-1} * b$ 
```

```
array([-1.,  1.])
```

We can think of solving a system as finding the linear combination of columns of  $A$  that equal  $b$

```
1 b_star = x[0] * A[:, 0] + x[1] * A[:, 1] # using x solution
2 print(f"b = {b}, b_star = {b_star}")
```

```
b = [2 1], b_star = [2. 1.]
```

## Column Space and Rank

- The column space of a matrix represents all possible linear combinations of its columns.
- It forms a basis for the space of solutions when solving systems of linear equations represented by the matrix
- The rank of a matrix is the dimension of its column space

```
1 A = np.array([[0, 2], [3, 4]])  
2 matrix_rank(A)
```

2

Hence, can solve  $Ax = b$  for any  $b \in \mathbb{R}^2$  since the column space is the entire space  $\mathbb{R}^2$

On the other hand, note

```
1 A = np.array([[1, 2],  
2               [2, 4]])  
3 matrix_rank(A)
```

1

So we can only solve  $Ax = b$  for  $b \propto \begin{bmatrix} 1 \\ 2 \end{bmatrix} \propto \begin{bmatrix} 2 \\ 4 \end{bmatrix}$

## Checking Singularity

```
1 A = np.array([[1, 2], [2, 4]])
2 # An (expensive) way to check if A is singular is if det(A) = 0
3 print(det(A) == 0.0)
4 print(matrix_rank(A) != A.shape[0]) # or check rank
5 # Check before inverting or use exceptions
6 try:
7     inv(A)
8     print("Matrix is not singular (invertible).")
9 except np.linalg.LinAlgError:
10    print("Matrix is singular (non-invertible).")
```

True

True

Matrix is singular (non-invertible).

## Determinant is Not Scale Invariant

- Reminder: numerical precision in calculations makes it hard to compare to zero
- The determinate is useful but depends on the scale of the matrix
- A more robust alternative is the condition number (more next lecture)

```
1 eps, K = 1e-8, 100000
2 A = np.array([[1, 2], [1 + eps, 2 + eps]])
3 print(f"det(A)={det(A):.5g}, det(K*A)={det(K*A):.5g}")
4 print(f"cond(A)={cond(A):.5g}, cond(K*A)={cond(A):.5g},")
5 print(f"det(inv(A))={det(inv(A)):.5g}, cond(inv(A))={cond(inv(A)):.5g}")
```

```
det(A)=-1e-08, det(K*A)=-100
cond(A)=1e+09, cond(K*A)=1e+09,
det(inv(A))=-1e+08, cond(inv(A))=1e+09
```

## Solving Linear Systems of Equations

---



## Solving Systems with Multiple RHS

- Inverse is nice because you can reuse the  $A^{-1}$  to solve  $Ax = b$  for many  $b$
- However, you can do this with `solve` as well
- Or can reuse LR factorizations (discussed next)

```
1 A = np.array([[0, 2], [3, 4]])
2 B = np.array([[2,3], [1,2]]) # [2,1] and [3,2] as columns
3 # or: B = np.column_stack([np.array([2, 1]),np.array([3,2])])
4 X = solve(A, B) # Solve AX = B for X
5 print(X)
6 print(f"Checking: A*{X[:,0]} = {A@X[:, 0]} = {B[:,0]}, column of B")
```

```
[[ -1.          -1.33333333]
 [  1.           1.5         ]]
```

```
Checking: A*[-1.  1.] = [2. 1.] = [2 1], column of B
```

## LU(P) Decompositions

- We can “factor” any square  $A$  into  $PA = LU$  for triangular  $L$  and  $U$ . Invertible can have  $A = LU$ , called the LU decomposition. “P” is for partial-pivoting
- Singular matrices may not have full-rank  $L$  or  $U$  matrices

```
1 A = np.array([[1, 2], [2, 4]])  
2 P, L, U = lu(A)  
3 print(f"L*U =\n{L @ U}")  
4 print(f"P*A =\n{P @ A}")
```

L\*U =

```
[[2. 4.]  
 [1. 2.]]
```

P\*A =

```
[[2. 4.]  
 [1. 2.]]
```

The  $P$  matrix is a permutation matrix of “pivots” the others are triangular

```
1 print(f"P =\n{P}")
2 print(f"L =\n{L}")
3 print(f"U =\n{U}")
```

P =

```
[[0. 1.]
 [1. 0.]]
```

L =

```
[[1.  0. ]
 [0.5 1. ]]
```

U =

```
[[2. 4.]
 [0. 0.]]
```

## LU Decompositions and Systems of Equations

- Pivoting is typically implied when talking about “LU”
- Used in the default `solve` algorithm (without more structure)
- Solving systems of equations with triangular matrices: for  $Ax = LUx = b$ 
  1. Define  $y = Ux$
  2. Solve  $Ly = b$  for  $y$  and  $Ux = y$  for  $x$
- Since both are triangular, process is  $O(N^2)$  (but LU itself  $O(N^3)$ )
- Could be used to find `inv`
  - $A = LU$  then  $AA^{-1} = I = LUA^{-1} = I$
  - Solve for  $Y$  in  $LY = I$ , then solve  $UA^{-1} = Y$
- Tight connection to textbook Gaussian elimination (including pivoting)

## LU for Non-Singular Matrices

```
1 A = np.array([[1, 2], [3, 4]])  
2 P, L, U = lu(A)  
3 print(f"L*U =\n{L @ U}")  
4 print(f"P*A =\n{P @ A}")
```

L\*U =

```
[[3. 4.]  
 [1. 2.]]
```

P\*A =

```
[[3. 4.]  
 [1. 2.]]
```

L, U, P

```
1 print(f"P =\n{P}")
2 print(f"L =\n{L}")
3 print(f"U =\n{U}")
```

P =

```
[[0. 1.]
 [1. 0.]]
```

L =

```
[[1.          0.          ]
 [0.33333333  1.          ]]
```

U =

```
[[3.          4.          ]
 [0.          0.66666667]]
```

## Backwards Substitution Example

$$Ux = b$$

$$U \equiv \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}, \quad b = \begin{bmatrix} 7 \\ 2 \end{bmatrix}$$

Solving bottom row for  $x_2$

$$2x_2 = 2, \quad x_2 = 1$$

Move up a row, solving for  $x_1$ , substituting for  $x_2$

$$3x_1 + 1x_2 = 7, \quad 3x_1 + 1 \times 1 = 7, \quad x_1 = 2$$

Generalizes to many rows. For  $L$  it is “forward substitution”

## Use Triangular Structure if Possible

- Triangular matrices of size  $N$  can be solved with back substitution in  $O(N^2)$
- Is  $O(N^2)$  good or bad? Beats,  $O(N^3)$  typical of general methods

```
1 U = np.array([[3, 1],  
2               [0, 2]])  
3 b = np.array([7, 2])  
4 solve(U, b) # works, but internally does an LU which is  $O(N^3)$   
5 solve_triangular(U, b, lower=False) # fast  $O(N^2)$ 
```

```
array([2., 1.])
```



## Symmetric Matrix Structure

Another common matrix type are symmetric,  $A = A^T$

```
1 A = np.array([[1, 2], [2, 5]]) # also posdef, not singular
2 b = np.array([1,4])
3 print(f"Symmetric A? {scipy.linalg.issymmetric(A)}")
4 solve(A, b, assume_a="sym") # could also use "pos" since positive definite
```

Symmetric A? True

array([-3., 2.])

## Positive Definite Matrices

- A symmetric matrix  $A$  is positive definite if  $x^T A x > 0$  for all  $x \neq 0$
- Useful in many areas, such as covariance matrices. Example

```
1 A = np.array([[1, 2], [2, 5]])  
2 x = np.array([0, 1]) # can't really check for all x  
3 print(f"x^T A x = {x.T @ A @ x}")
```

$$x^T A x = 5$$

- Example of a symmetric matrix that is not positive definite

```
1 A = np.array([[1, 2], [2, 0]])  
2 print(f"x^T A x = {x.T @ A @ x}") # one counterexample is enough
```

$$x^T A x = 0$$

- We can check these with eigenvalues

# Cholesky Decomposition

- For symmetric positive definite matrices:  $L = U^T$ ,
- Called a Cholesky decomposition:  $A = LL^T$  for a lower triangular matrix  $L$

```
1 A = np.array([[1, 2], [2, 5]])
2 L = cholesky(A)
3 print(L)
4 print(f"L*L^T =\n{L @ L.T}")
```

```
[[1. 2.]
 [0. 1.]]
L*L^T =
[[5. 2.]
 [2. 1.]]
```

## Solving Positive Definite Systems

```
1 A = np.array([[1, 2], [2, 5]])
2 b = np.array([1, 4])
3 print(solve(A, b, assume_a="pos")) # uses cholesky internally
4
5 L = cholesky(A)
6 y = solve_triangular(L, b, lower=True)
7 x = solve_triangular(L.T, y, lower=False)
8 print(x)
```

[-3. 2.]

[1. 4.]

- Covariance matrices are positive-definite, semi-definite if degenerate
- Key property of Gaussian random variables:
  - $X \sim N(\mu, \Sigma)$  for  $\mu \in \mathbb{R}^N, \Sigma \in \mathbb{R}^{N \times N}$
  - $X = \mu + AZ$  for  $Z \sim N(0_N, I_N)$  where  $AA^T = \Sigma$
- That is,  $A$  is the Cholesky decomposition of the covariance matrix

# Eigenvalues and Eigenvectors

---

# Eigenvalues and Eigenvectors

- For a square  $A$ , an eigenvector  $x$  and eigenvalue  $\lambda$  satisfy

$$Ax = \lambda x$$

- $A \in \mathbb{R}^{N \times N}$  has  $N$  eigenvalue/eigenvector pairs, possible multiplicity of  $\lambda$
- Intuition:  $x$  is a direction  $Ax \propto x$  and  $\lambda$  says how much it “stretches”
- Properties:
  - For any eigenvector  $x$  and scalar  $c$  then  $cx \propto Ax$  as well
  - Symmetric matrices have real eigenvalues and orthogonal eigenvectors. i.e.  $x_1 \cdot x_2 = 0$  for  $x_1 \neq x_2$  eigenvectors. Complex in general
  - Singular if and only if it has an eigenvalue of zero
  - Positive (semi)definite if and only if all eigenvalues are strictly (weakly) positive
  - Diagonal matrix has eigenvalues as its diagonal
  - Triangular matrix has eigenvalues as its diagonal

## Positive Definite and Eigenvalues

You cannot check  $x^T Ax > 0$  for all  $x$ . Check if “stretching” is positive

```
1 A = np.array([[3, 1], [2, 1]])
2 # A_eigs = np.real(eigvals(A)) # symmetric matrices have real eigenvalues
3 A_eigs = eigvalsh(A) # specialized for symmetric/hermitian matrices
4 print(A_eigs)
5 is_positive_definite = np.all(A_eigs > 0)
6 is_positive_semi_definite = np.all(A_eigs >= 0) # or eigvals(A) >= -eps
7 print(f"pos-def? {is_positive_definite}")
8 print(f"pos-semi-def? {is_positive_semi_definite}")
```

```
[-0.23606798  4.23606798]
```

```
pos-def? False
```

```
pos-semi-def? False
```



## Positive Semi-Definite Matrices May Have a Zero Eigenvalue

The simplest positive-semi-definite (but not posdef) matrix is

```
1 A_eigs = eigvalsh(np.array([[1, 0], [0, 0]]))
2 print(A_eigs)
3 is_positive_definite = np.all(A_eigs > 0)
4 is_positive_semi_definite = np.all(A_eigs >= 0) # or eigvals(A) >= -eps
5 print(f"pos-def? {is_positive_definite}")
6 print(f"pos-semi-def? {is_positive_semi_definite}")
```

[0. 1.]

pos-def? False

pos-semi-def? True

# Eigenvalue Decomposition

- For square, symmetric, non-singular matrix  $A$  factor into

$$A = Q\Lambda Q^{-1}$$

- $Q$  is a matrix of eigenvectors,  $\Lambda$  is a diagonal matrix of eigenvalues (in order)
- For symmetric matrices, the eigenvectors are orthogonal and  $Q^{-1}Q = Q^T Q = I$ , and we can think of them as forming an orthonormal basis
- Orthogonal matrices can be thought of as rotations without stretching
- More general matrices all have a Singular Value Decomposition (SVD)
- With symmetric  $A$ , an interpretation of  $Ax$  is that we can first rotate  $x$  into the  $Q$  basis, then stretch by  $\Lambda$ , then rotate back
- Can be used to find  $A^t$  for large  $t$  (e.g. for Markov chains)
  - $P^t$ , i.e.  $P \cdot P \cdot \dots \cdot P$  for  $t$  times
  - $P = Q\Lambda Q^{-1}$  then  $P^t = Q\Lambda^t Q^{-1}$  where  $\Lambda^t$  is just the pointwise power

## Eigenvalue Decomposition of Symmetric Matrix Example

```
1 A = np.array([[2, 1], [1, 3]])
2 Lambda, Q = eig(A)
3 print(f"eigenvectors are column-by-column in Q =\n{Q}")
4 print(f"eigenvalues are in Lambda = {Lambda}")
5 print(f"Q Lambda Q^T =\n{Q @ np.diag(np.real(Lambda)) @ Q.T}")
```

eigenvectors are column-by-column in Q =

```
[[-0.85065081 -0.52573111]
 [ 0.52573111 -0.85065081]]
```

eigenvalues are in Lambda = [1.38196601+0.j 3.61803399+0.j]

Q Lambda Q^T =

```
[[2. 1.]
 [1. 3.]]
```

## Spectral Radius is Maximum Absolute Eigenvalue

- If any  $\lambda \in \Lambda$  are  $> 1$  can see this would explode
- Useful for seeing if iteration  $x_{t+1} = Ax_t$  from a  $x_0$  explodes

### Definition (Spectral Radius)

The spectral radius of matrix  $A$  is  $\rho(A) = \max_{\lambda \in \Lambda} |\lambda|$

## Least Squares and the Normal Equations

---

Given a matrix  $X \in \mathbb{R}^{N \times M}$  and a vector  $y \in \mathbb{R}^N$ , we want to find  $\beta \in \mathbb{R}^M$  such that

$$\min_{\beta} \|y - X\beta\|^2, \text{ that is,}$$
$$\min_{\beta} \sum_{n=1}^N \frac{1}{N} (y_n - X_n \cdot \beta)^2$$

Where  $X_n$  is n'th row. Take FOCS and rearrange to get

$$(X^T X)\beta = X^T y$$

## Solving the Normal Equations

- The  $X$  is often referred to as the “design matrix”.  $X^T X$  as the Gram matrix
- Can form  $A = X^T X$  and  $b = X^T y$  and solve  $A\beta = b$ .
  - Or invert  $X^T X$  to get  $\beta = (X^T X)^{-1} X^T y$
  - Note that  $X^T X$  is symmetric and, if  $X$  is full-rank, positive definite
- In practice, use the `lstsq` function in `scipy`
  - It uses better algorithms using eigenvectors. More stable (see next lecture on conditioning)
  - One algorithm uses another factoring, the QR decomposition
  - There,  $X = QR$  for  $Q$  orthogonal and  $R$  upper triangular. See QR Decomposition for more
- We are exploring linear algebra in this lecture
  - For applied work use higher-level libraries like `statsmodels` (integrated well with `pandas` and `seaborn`)
  - See `statsmodels` docs for R-style notation
  - See QuantEcon OLS for later

## Example of LLS using Scipy

```
1 N, M = 100, 5
2 X = np.random.randn(N, M)
3 beta = np.random.randn(M)
4 y = X @ beta + 0.05 * np.random.randn(N)
5 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X, y)
6 print(f"beta =\n {beta}\nbeta_hat =\n{beta_hat}")
```

beta =

```
[-0.6745072  2.0103398  0.89741915  0.62875708  0.64840828]
```

beta\_hat =

```
[-0.66550515  2.00553235  0.89398807  0.624045  0.65692653]
```



## Solving using the Normal Equations

Or we can solve it directly. Provide matrix structure (so it can use a Cholesky)

```
1 beta_hat = solve(X.T @ X, X.T @ y, assume_a="pos")
2 print(f"beta =\n {beta}\nbeta_hat =\n{beta_hat}")
```

```
beta =
 [-0.6745072  2.0103398  0.89741915  0.62875708  0.64840828]
beta_hat =
 [-0.66550515  2.00553235  0.89398807  0.624045  0.65692653]
```

## Collinearity in “Tall” Matrices

- Tall  $\mathbb{R}^{N \times M}$  “design matrices” have  $N > M$  and are “overdetermined”
- The rank of a matrix is full rank if all columns are linearly independent
- You can only identify  $M$  parameters with  $M$  linearly independent columns

```
1 X = np.array([[1, 2], [2, 5], [3, 7]]) # 3 observations, 2 variables
2 X_col = np.array([[1, 2], [2, 4], [3, 6]]) # all proportional
3 print(f"rank(X) = {matrix_rank(X)}, rank(X_col) = {matrix_rank(X_col)}")
```

```
rank(X) = 2, rank(X_col) = 1
```

- If  $X$  is not full rank, then  $X^T X$  is not invertible. For example:

```
1 print(f"cond(X'*X)={cond(X.T@X)}, cond(X_col'*X_col)={cond(X_col.T@X_col)}")
```

```
cond(X'*X)=2819.332978639814, cond(X_col'*X_col)=1.1014450683078442e+16
```

- Note that when you start doing operations on matrices, numerical error creeps in, so you will not get an exact number
- The rule-of-thumb with condition numbers is that if it is  $1 \times 10^k$  then you lose about  $k$  digits of precision. So this effectively means it is singular
- Given the singular matrix, this means a continuum of  $\beta$  will solve the problem

## lstsq Solves it? Careful on Interpretation!

- Since  $X_{col}^T X_{col}$  is singular, we cannot use `solve(X.T@X, y)`
- But what about `lstsq` methods?
- As you will see, this gives an answer. Interpretation is hard
- The key is that in the case of non-full rank, you cannot identify individual parameters
  - Related to “Identification” in econometrics
  - Having low residuals is not enough

```
1 y = np.array([5.0, 10.1, 14.9])
2 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X_col, y)
3 print(f"beta_hat_col = {beta_hat}")
4 print(f"rank={rank}, cols={X.shape[1]}, norm(X*beta_hat_col-y)={norm(residuals)}
```

```
beta_hat_col = [0.99857143 1.99714286]
rank=1, cols=2, norm(X*beta_hat_col-y)=0.0
```

## Fat Design Matrices

- Fat  $\mathbb{R}^{N \times M}$  “design matrices” have  $N < M$  and are “underdetermined”
- Less common in econometrics, but useful to understand the structure
- A continuum  $\beta \in \mathbb{R}^{M - \text{rank}(X)}$  solve this problem

```
1 X = np.array([[1, 2, 3], [0, 5, 7]]) # 2 rows, 3 variables
2 y = np.array([5, 10])
3 beta_hat, residuals, rank, s = scipy.linalg.lstsq(X, y)
4 print(f"beta_hat = {beta_hat}, rank={rank}, ? residuals = {residuals}")
```

```
beta_hat = [0.8 0.6 1. ], rank=2, ? residuals = []
```

## Which Solution?

- Residuals are zero here because there are enough parameters to fit perfectly (i.e., it is underdetermined)
- Given the multiple solutions, the `lstsq` is giving

$$\min_{\beta} \|\beta\|_2^2 \text{ s.t. } X\beta = y$$

- i.e., the “smallest” coefficients which interpolate the data exactly
- Which trivially fulfills the OLS objective:  $\min_{\beta} \|y - X\beta\|_2^2$
- Useful and common in ML, but be **very** careful when interpreting for economics
  - Tight connections to Bayesian versions of statistical tests
  - But until you understand econometrics and “identification” well, **stick to full-rank matrices**
  - **Advanced topics:** search for “Regularization”, “Ridgeless Regression” and “Benign Overfitting in Linear Regression.”